

# Sustainable Computing

Jan Rybizki (IAB Weimar, formerly MPIA Heidelberg)

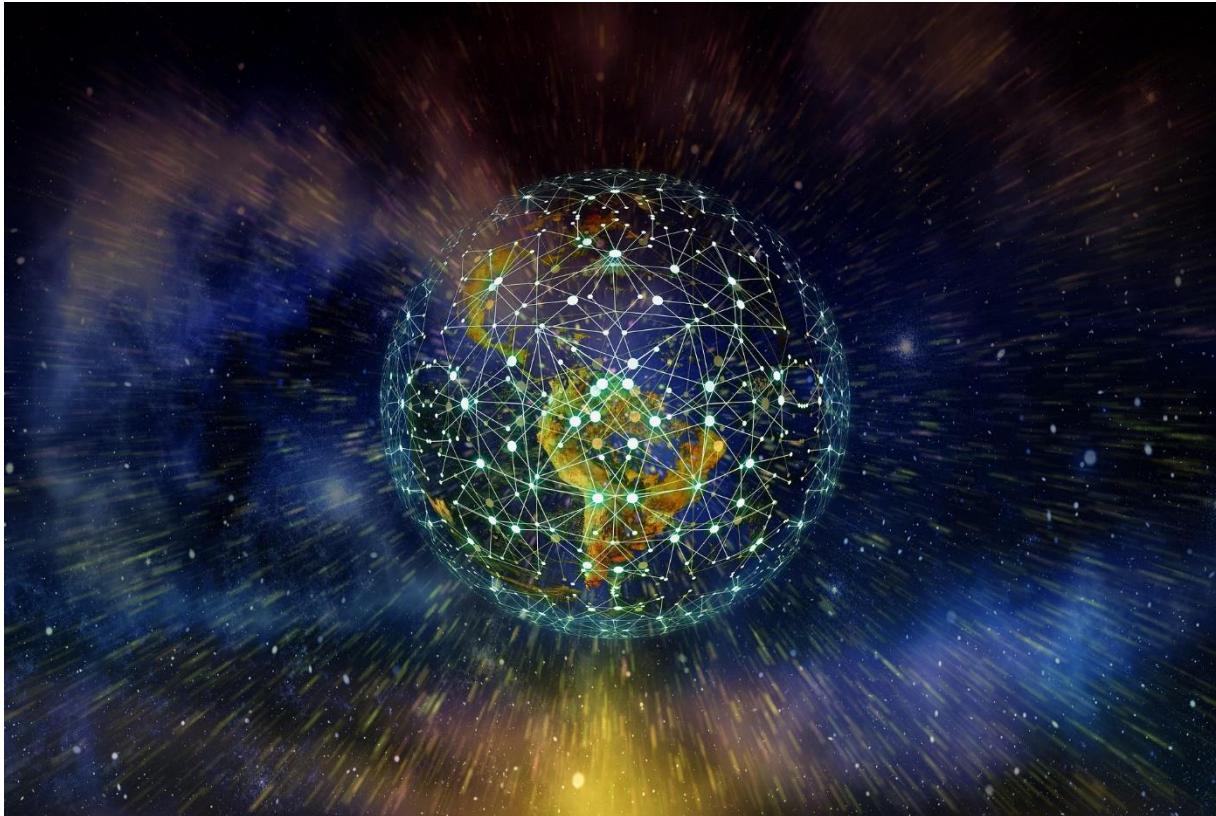
[j.rybizki@iab-weimar.de](mailto:j.rybizki@iab-weimar.de)

Twitter: @rybizki

- What do you think is sustainable computing?
- Please vote via handsraising



# Climate modelling



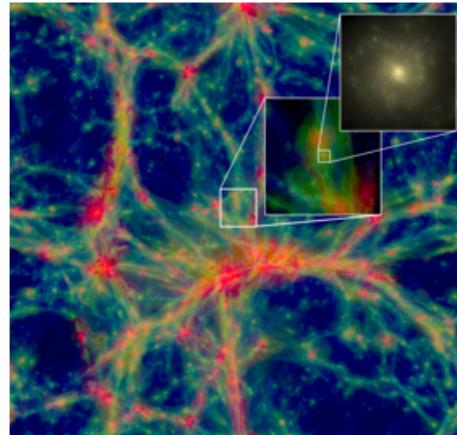
## Super simulations help determine the mass of neutrinos

06/23/2020

Felix Spanier from the Center for Astronomy of the University (ZAH) and his collaboration partners Guido Drexlin and Jonas Kellerer from the Karlsruhe Institute of Technology (KIT) have achieved great success in the current application round of the Gauss Center for Supercomputing (GCS): 35 million CPU hours can be used for plasma simulations as part of the Karlsruhe Tritium Neutrino Experiment (KATRIN) whose aim is to measure the mass of the neutrino. The result of this measurement is of fundamental importance, among other things, for elementary particle physics or our understanding of the development of structures in the universe.

The KATRIN experiment is a scientific superlative, the construction of which took about 15 years in international cooperation. KATRIN exploits the so-called beta decay of radioactive tritium gas. This decay produces an electron and neutrino at the same time. The neutrino cannot be detected. But KATRIN can measure a tiny anomaly in the energy distribution of decay electrons and conclude the mass of the neutrino.

However, the tiny anomaly in the energy spectrum is also influenced by the behaviour of the tritium gas. The decay of the tritium frees up high-energy electrons that can ionize the gas itself, so that it becomes a plasma with special electromagnetic properties. The effect on the electrons to be measured is actually minimal, but in order to measure the mass of the neutrino to the required accuracy of 0.000 000 000 000 000 000 000 000 000 3 g, the influence of the plasma must be known indeed very precisely.

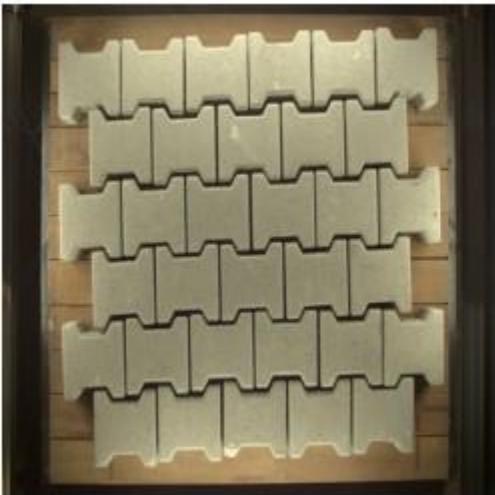


Visualization of a simulation of the formation of cosmic structures. As the most common elementary particle in the universe, the neutrino has a significant impact. Therefore, it is important to measure its mass very accurately. (Copyright: The EAGLE-Project <http://icc.dur.ac.uk/Eagle/>)

# Quality control in concrete production

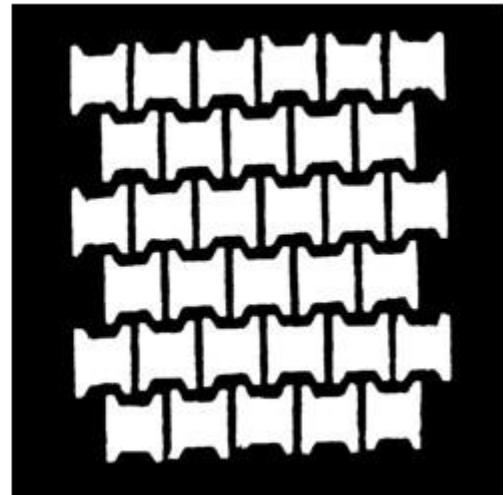


Input image

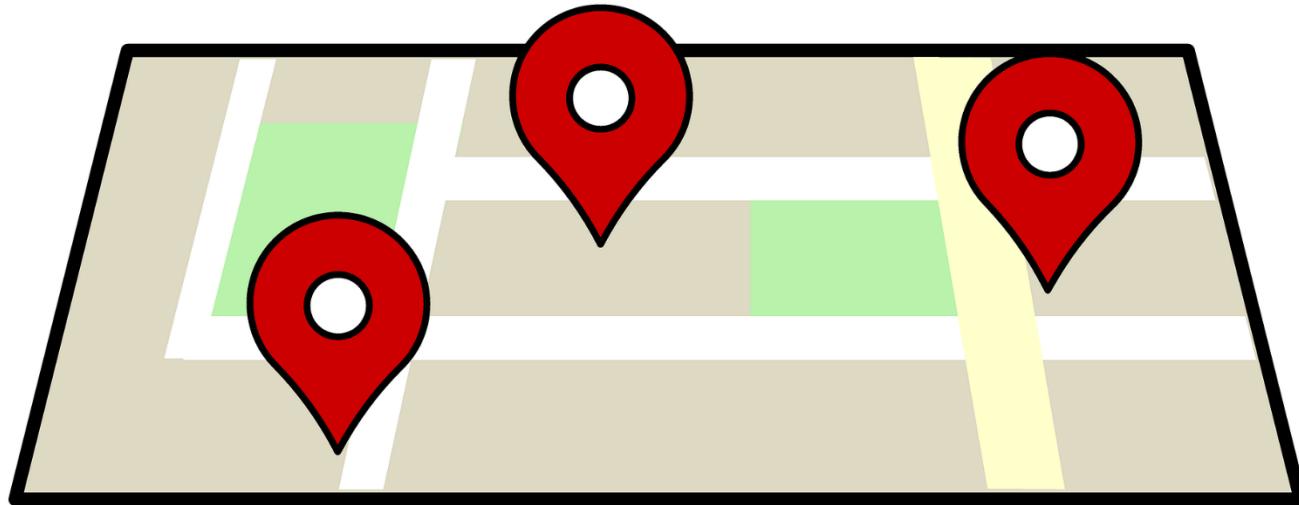


segmentation

Paving stone mask



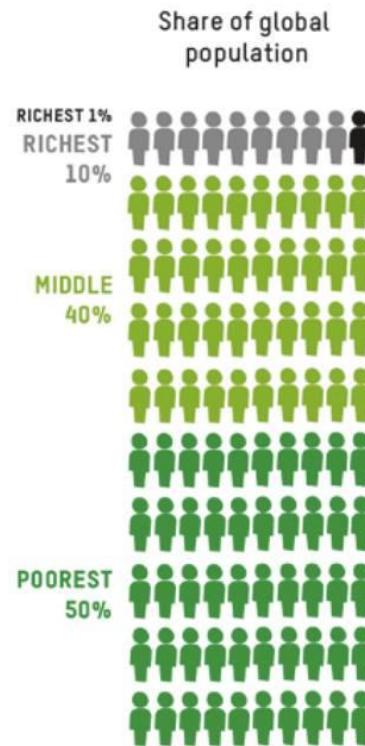
# Navigation



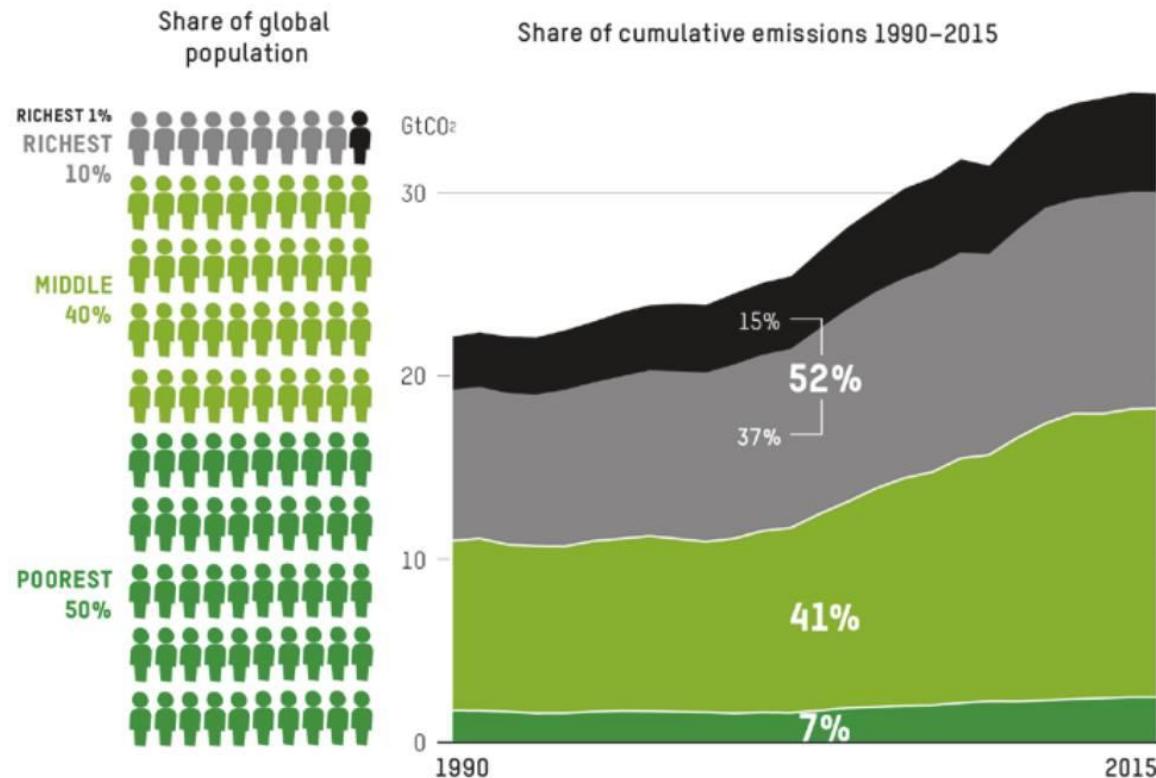
Anthropogenic greenhouse gas emissions are problematic



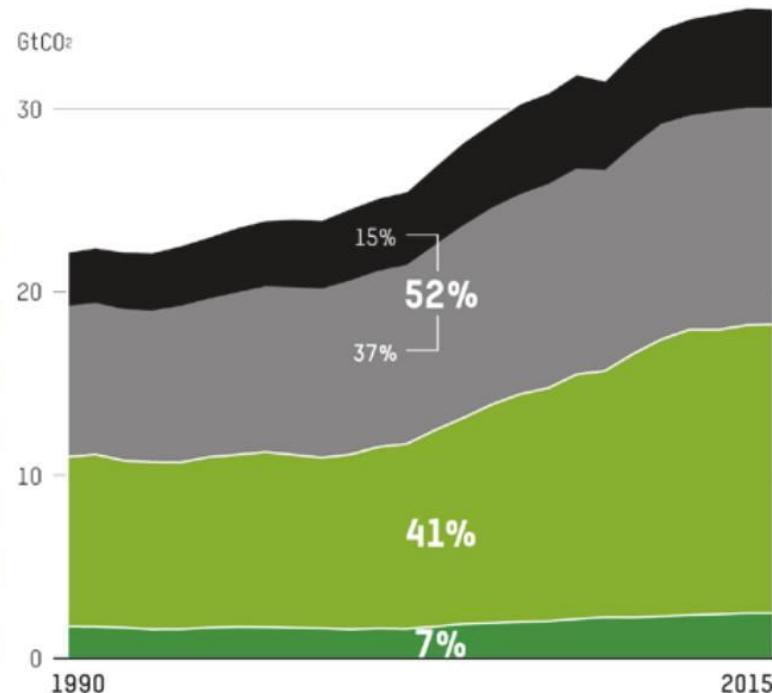
# Introduction



# Introduction



# Introduction



# Introduction

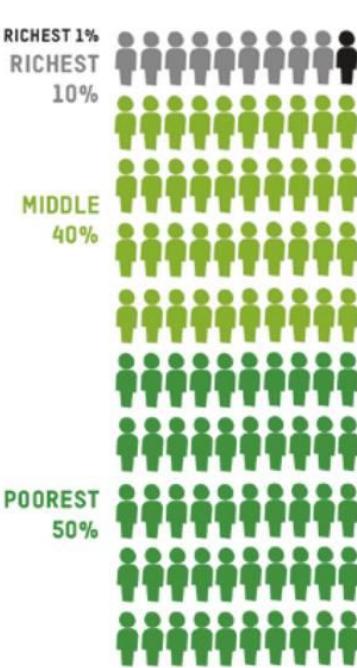
average emission

75t CO<sub>2eq</sub>/ yr  
23t CO<sub>2eq</sub>/ yr

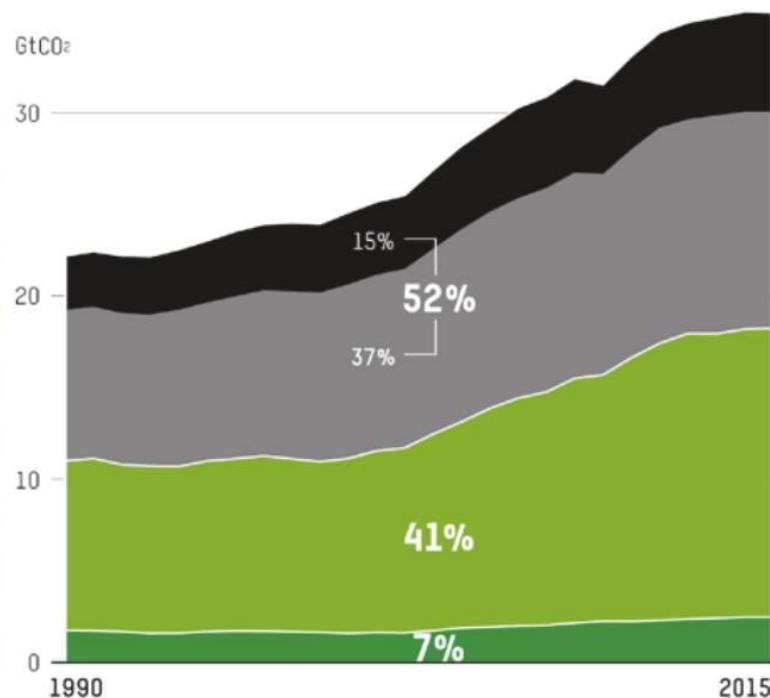
income threshold

> 100k € / yr  
> 35k € / yr

Share of global population



Share of cumulative emissions 1990–2015



5t CO<sub>2eq</sub>/ yr

> 5k € / yr

0.75t CO<sub>2eq</sub>/ yr

< 5k € / yr

# Introduction

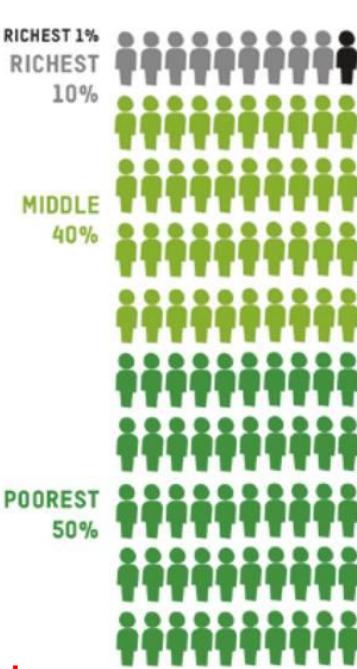
average emission

75t CO<sub>2eq</sub>/ yr  
23t CO<sub>2eq</sub>/ yr

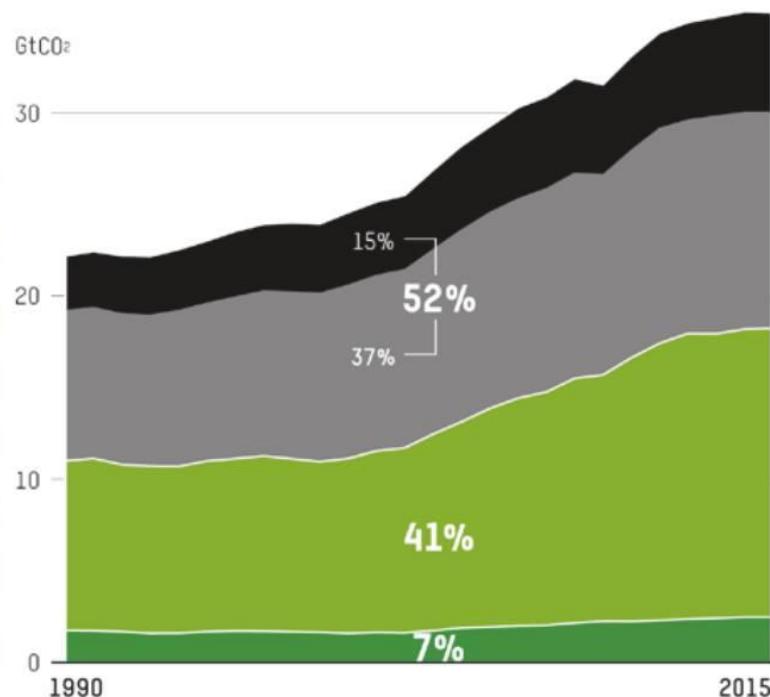
income threshold

> 100k € / yr  
> 35k € / yr

Share of global population



Share of cumulative emissions 1990–2015



2t CO<sub>2eq</sub>/ yr is sustainable!

Glad to speak to you!

- GHG study of MPIA and Australian Astro community
- Background
  - Electricity grid in Germany
  - Computer physical boundaries
  - Data centers energy usage
- Increasing computing efficiency
- Solution ideas

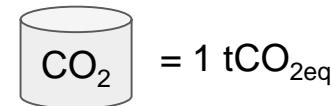
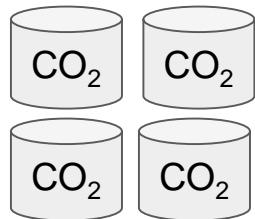
**annual tCO<sub>2eq</sub> per MPIA researcher**

# annual tCO<sub>2eq</sub> per MPIA researcher

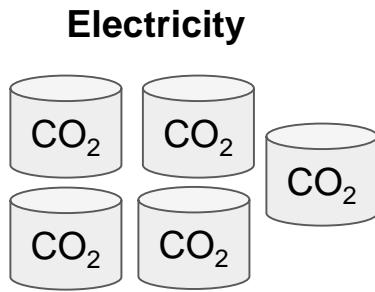
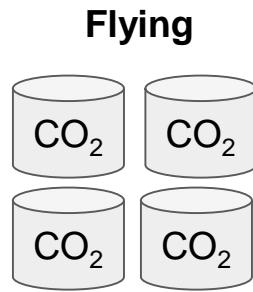
data taken from [Jahnke+ 2020](#) and refined by reanalysis

# annual tCO<sub>2eq</sub> per MPIA researcher

Flying



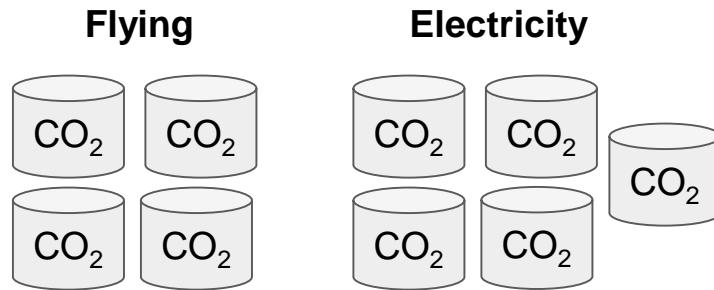
# annual tCO<sub>2eq</sub> per MPIA researcher



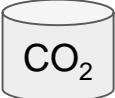
= 1 tCO<sub>2eq</sub>

A single cylinder, labeled  $\text{CO}_2$ , representing the equivalent of one tonne of CO<sub>2</sub> equivalents.

# annual tCO<sub>2eq</sub> per MPIA researcher

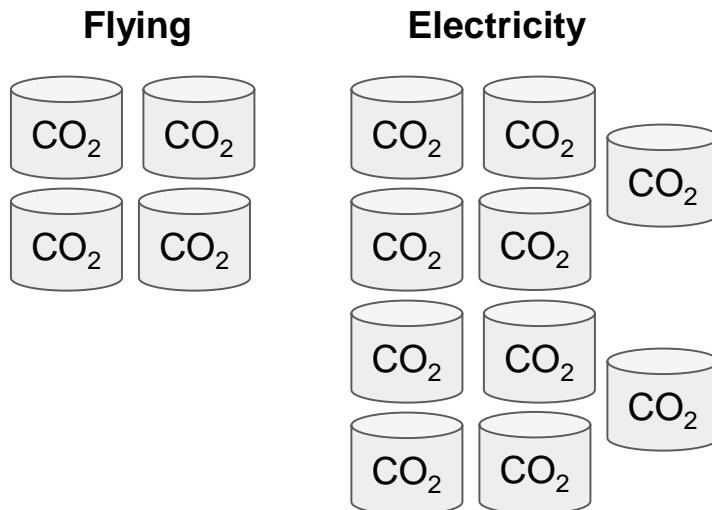


Supercomputing facilities  
make up ~80% of electricity  
usage



= 1 tCO<sub>2eq</sub>

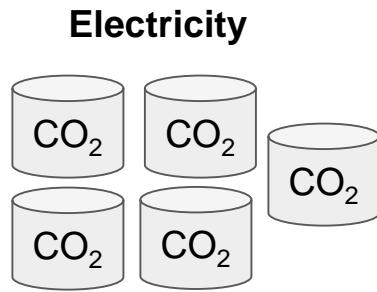
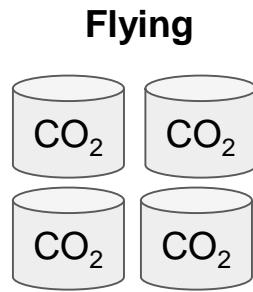
# annual tCO<sub>2eq</sub> per MPIA researcher



Emissions if we would use the  
CO<sub>2eq</sub>-intensity of the **German**  
electricity production instead of  
**Heidelberg / Munich** city grid

= 1 tCO<sub>2eq</sub>

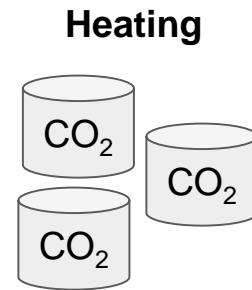
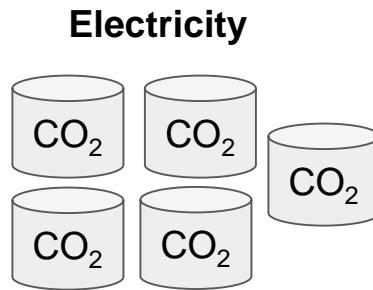
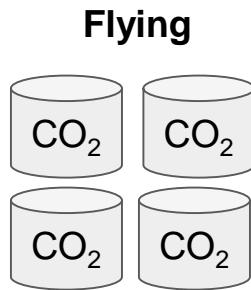
# annual tCO<sub>2eq</sub> per MPIA researcher



= 1 tCO<sub>2eq</sub>

A single cylinder, labeled  $\text{CO}_2$ , representing the equivalent of one tonne of CO<sub>2</sub> equivalent greenhouse gas emissions.

# annual tCO<sub>2eq</sub> per MPIA researcher



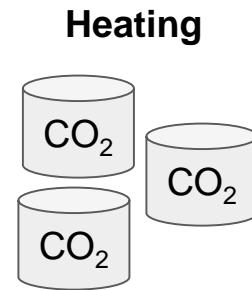
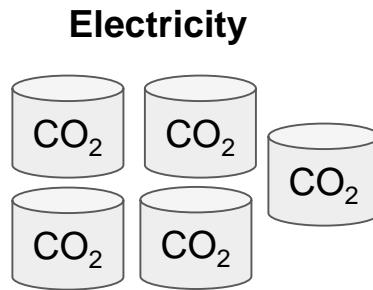
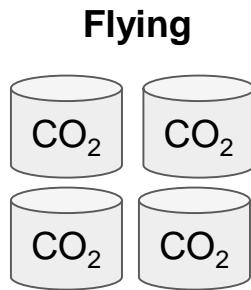
= 1 tCO<sub>2eq</sub>

A single cylinder, labeled  $\text{CO}_2$ , representing the equivalent of one tonne of CO<sub>2</sub> equivalents.



Carolin Liefke, HdA

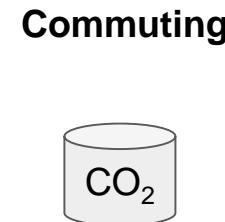
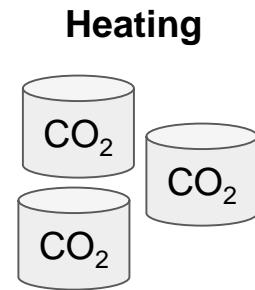
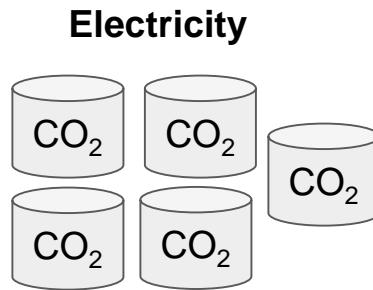
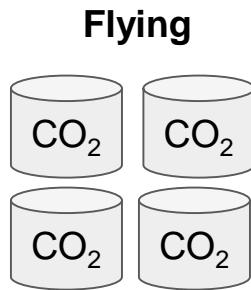
# annual tCO<sub>2eq</sub> per MPIA researcher



= 1 tCO<sub>2eq</sub>

A single cylinder, labeled  $\text{CO}_2$ , representing the equivalent of one tonne of CO<sub>2</sub> equivalents.

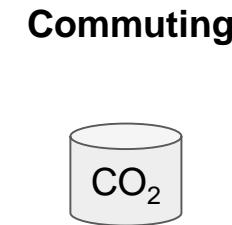
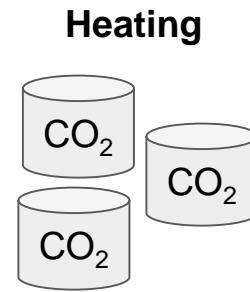
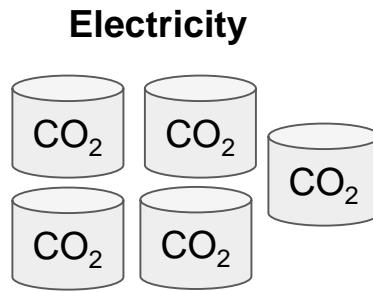
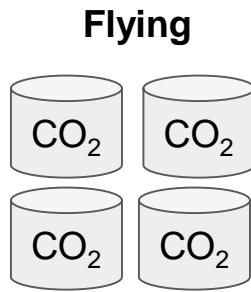
# annual tCO<sub>2eq</sub> per MPIA researcher



= 1 tCO<sub>2eq</sub>

A single cylinder, labeled  $\text{CO}_2$ , representing the unit of measurement for the annual tCO<sub>2eq</sub> per MPIA researcher.

# annual tCO<sub>2eq</sub> per MPIA researcher

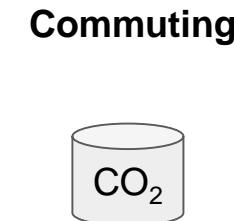
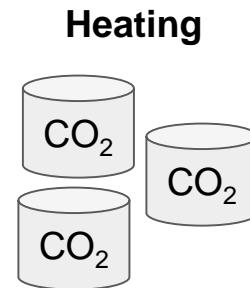
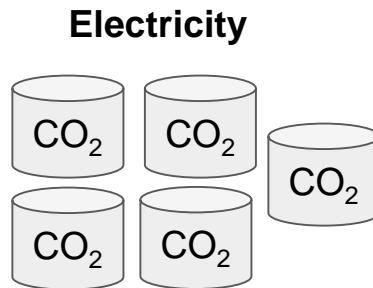
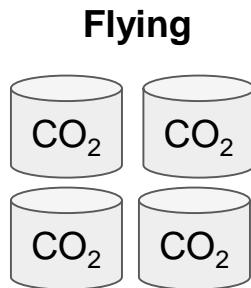


?

**Infrastructure**

CO<sub>2</sub> = 1 tCO<sub>2eq</sub>

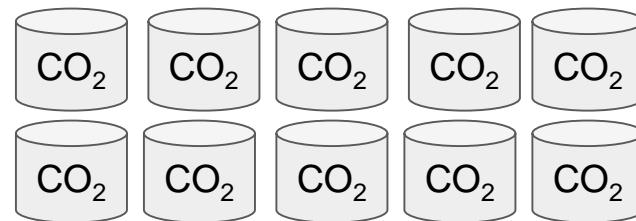
# annual tCO<sub>2eq</sub> per MPIA researcher



?

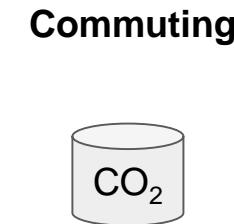
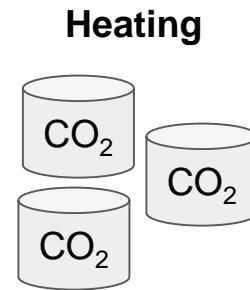
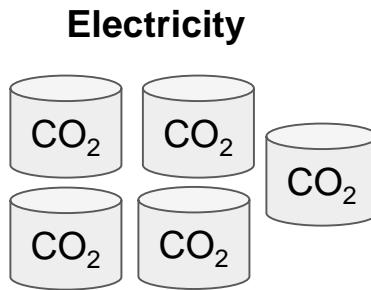
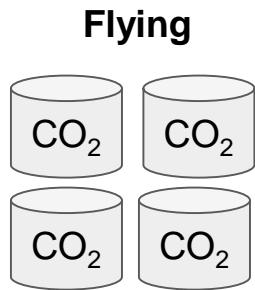
**Infrastructure**

**German yearly average**



CO<sub>2</sub> = 1 tCO<sub>2eq</sub>

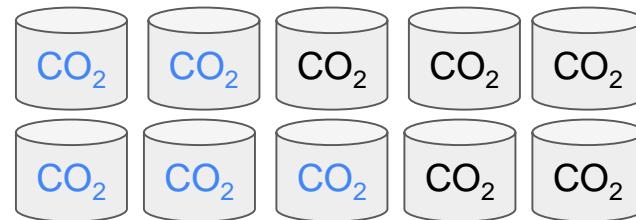
# annual tCO<sub>2eq</sub> per MPIA researcher



?

**Infrastructure**

**German yearly average**



**German Pledge 2030**

CO<sub>2</sub> = 1 tCO<sub>2eq</sub>

# Supercomputing put into perspective

# Supercomputing put into perspective

- Average German uses

**5.9 MWh per year = 1.5MWh (private) + 4.4MWh (work)** of electricity

# Supercomputing put into perspective

- Average German uses  
**5.9 MWh per year** = 1.5MWh (private) + 4.4MWh (work) of electricity
- Per Astronomer (MPIA and Australian study), we use around  
**20 MWh per year** for supercomputing

# Supercomputing put into perspective

- Average German uses  
**5.9 MWh per year** = 1.5MWh (private) + 4.4MWh (work) of electricity
- Per Astronomer (MPIA and Australian study), we use around  
**20 MWh per year** for supercomputing

2.3 kW (200 cores)

**non-stop**

# Supercomputing put into perspective

- Average German uses  
**5.9 MWh per year** = 1.5MWh (private) + 4.4MWh (work) of electricity
- Per Astronomer (MPIA and Australian study), we use around  
**20 MWh per year** for supercomputing

2.3 kW (200 cores)

non-stop

=



(2 vacuum cleaners)

# Supercomputing put into perspective

- Average German uses  
**5.9 MWh per year** = 1.5MWh (private) + 4.4MWh (work) of electricity
- Per Astronomer (MPIA and Australian study), we use around  
**20 MWh per year** for supercomputing

2.3 kW (200 cores)

non-stop

=



(2 vacuum cleaners)  
all year long...

# Hardware life-cycle\*

\*mostly production and shipping  
vs. consumption

# Hardware life-cycle\*

Hardware is used 5 years, after which maintenance gets really expensive

\*mostly production and shipping  
vs. consumption

# Hardware life-cycle\*

Hardware is used 5 years, after which maintenance gets really expensive

Production and shipping add 20-30% CO<sub>2eq</sub> costs (at 100% utilization)

\*mostly production and shipping  
vs. consumption

# Hardware life-cycle\*

Hardware is used 5 years, after which maintenance gets really expensive

Production and shipping add 20-30% CO<sub>2eq</sub> costs (at 100% utilization)

For office equipment it can be 50% or more

\*mostly production and shipping  
vs. consumption

# Hardware life-cycle\*

Hardware is used 5 years, after which maintenance gets really expensive

Production and shipping add 20-30% CO<sub>2eq</sub> costs (at 100% utilization)

For office equipment it can be 50% or more

Check: <https://ecoinfo.cnrs.fr/ecoddiag-calcul/> to get life-cycle estimates

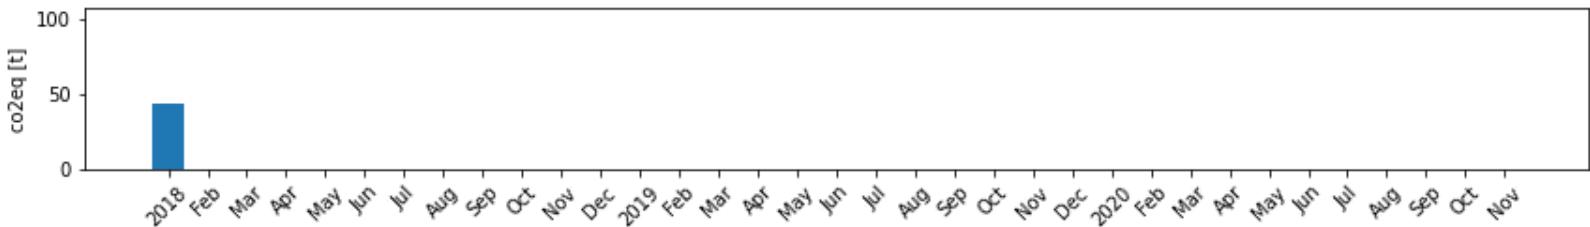
\*mostly production and shipping  
vs. consumption

# An assessment of MPIA's travel activities

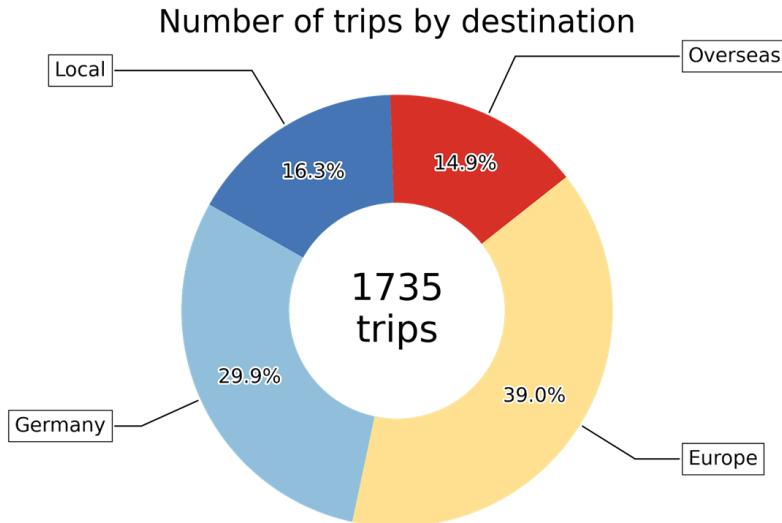


2018-2020

[https://github.com/jan-rybizki/Business\\_trips\\_carbon\\_footprint](https://github.com/jan-rybizki/Business_trips_carbon_footprint)

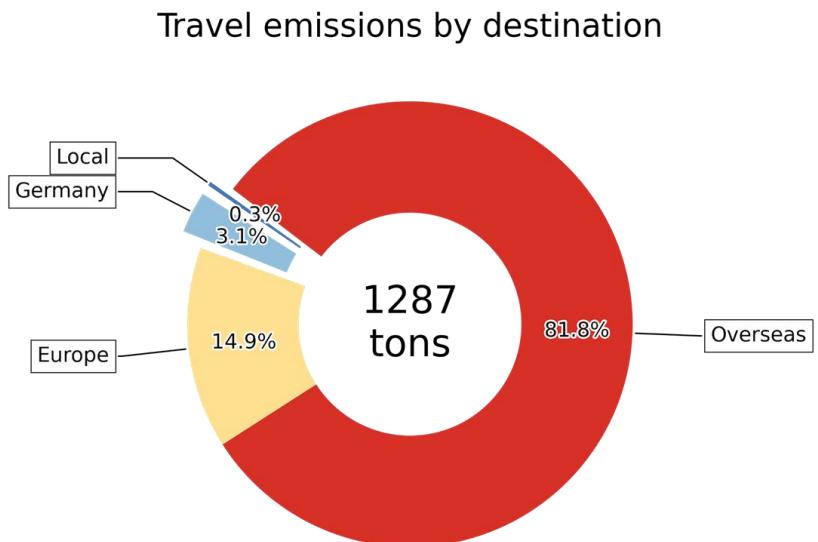
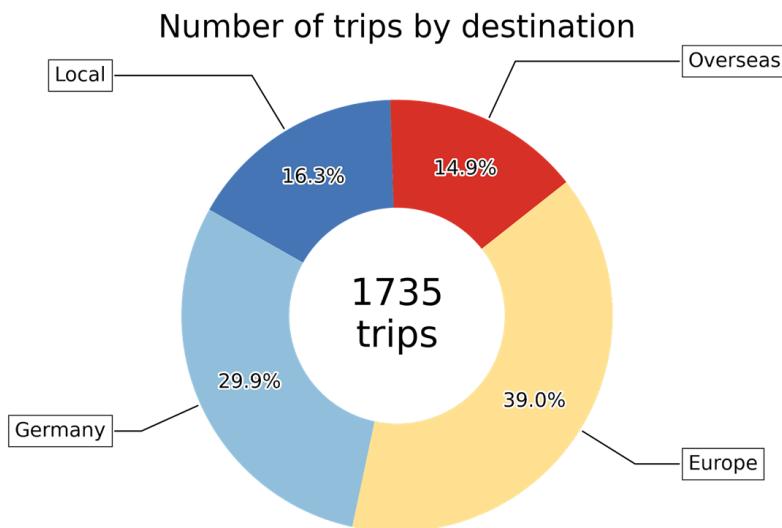


# Number of flights vs. CO<sub>2eq</sub> emissions by destination



Overseas make up **15%** of trips **by number...**

# Number of flights vs. CO<sub>2eq</sub> emissions by destination



Overseas make up **15%** of trips **by number...**

**...but 82% of trips by CO<sub>2eq</sub> emissions.**

# Online meetings help!



- Less emissions (a factor of 1000 less)
- More inclusive

# Background knowledge



- Electricity grid in Germany
- Computer physical boundaries
- Data centers energy usage

# **Electric supply - boundary conditions in the near future**

## **Electric supply - boundary conditions in the near future**

- heating and mobility will get electrified → more demand

## **Electric supply - boundary conditions in the near future**

- heating and mobility will get electrified → more demand
- nuclear and coal and gas need to be reduced → less supply

## **Electric supply - boundary conditions in the near future**

- heating and mobility will get electrified → more demand
- nuclear and coal and gas need to be reduced → less supply
- CO<sub>2eq</sub> price tag will increase → electricity will be more expensive

## Electric supply - boundary conditions in the near future

- heating and mobility will get electrified → more demand
- nuclear and coal and gas need to be reduced → less supply
- CO<sub>2eq</sub> price tag will increase → electricity will be more expensive
- share of renewables will increase → volatile supply

# Germany + electricity consumption

# Germany + electricity consumption

**Question 1:** how much electricity  
does Germany consume?

# Germany + electricity consumption

**Question 1:** how much electricity does Germany consume?

→ 488 TWh (2019)

→ 5.9 MWh per person per year

# Germany + electricity consumption

**Question 1:** how much electricity does Germany consume?

→ 488 TWh (2019)

→ 5.9 MWh per person per year

**Question 2:** how much CO<sub>2eq</sub> does this emit?

# Germany + electricity consumption

**Question 1:** how much electricity does Germany consume?

→ 488 TWh (2019)

→ 5.9 MWh per person per year

**Question 2:** how much CO<sub>2eq</sub> does this emit?

→ ?? kg CO<sub>2eq</sub> / MWh

# Germany + electricity consumption

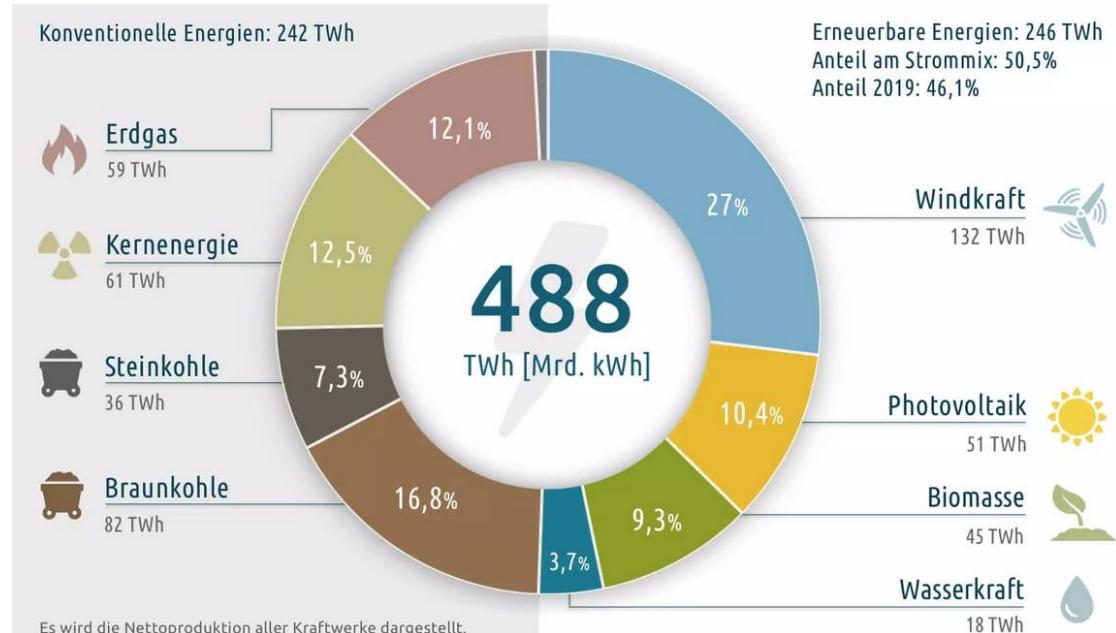
**Question 1:** how much electricity does Germany consume?  
a look into Germany's **electricity mix...**

→ 488 TWh (2019)

→ 5.9 MWh per person per year

**Question 2:** how much CO<sub>2eq</sub> does this emit?

→ ?? kg CO<sub>2eq</sub> / MWh



\*each country has a different mix of electricity sources

**CO<sub>2</sub>eq-Emissions German electricity mix [kg / MWh]**

Quelle: EUPD Research 2021

	Jan	Feb	Mrz	Apr	Mai	Jun	Jul	Aug	Sep	Okt	Nov	Dez
00:00	349	227	297	294	314	362	408	413	450	368	433	396
01:00	346	222	293	287	305	348	399	401	443	362	429	390
02:00	342	217	293	284	301	341	390	393	438	352	425	387
03:00	341	216	290	284	302	340	385	393	437	359	426	385
04:00	343	220	297	290	306	344	386	399	442	368	430	388
05:00	355	232	310	301	319	357	393	414	455	389	441	400
06:00	367	247	322	304	318	360	389	421	464	407	451	414
07:00	373	253	311	288	295	342	360	403	449	407	450	418
08:00	374	253	290	263	268	318	327	373	419	400	440	419
09:00	367	247	273	235	242	294	296	346	393	386	430	416
10:00	359	239	253	203	216	274	273	322	366	372	420	410
11:00	354	232	240	183	200	259	257	302	341	361	412	407
12:00	350	229	235	172	192	250	247	291	324	351	412	409
13:00	353	228	238	167	189	245	242	288	315	349	424	419
14:00	364	233	246	168	189	245	242	292	320	356	444	431
15:00	380	244	263	176	197	251	250	306	340	374	469	437
16:00	389	258	292	198	210	265	269	328	377	398	476	429
17:00	379	264	323	239	240	288	300	360	420	413	460	417
18:00	373	260	335	279	276	320	333	390	454	411	454	418
19:00	372	259	333	301	305	352	364	414	462	404	456	421
20:00	368	254	330	306	324	376	393	427	462	400	460	422
21:00	363	246	326	307	332	388	409	431	468	394	456	418
22:00	359	243	323	305	333	386	413	432	468	386	454	416
23:00	351	235	316	296	326	383	417	432	457	373	448	409

# Germany + electricity consumption

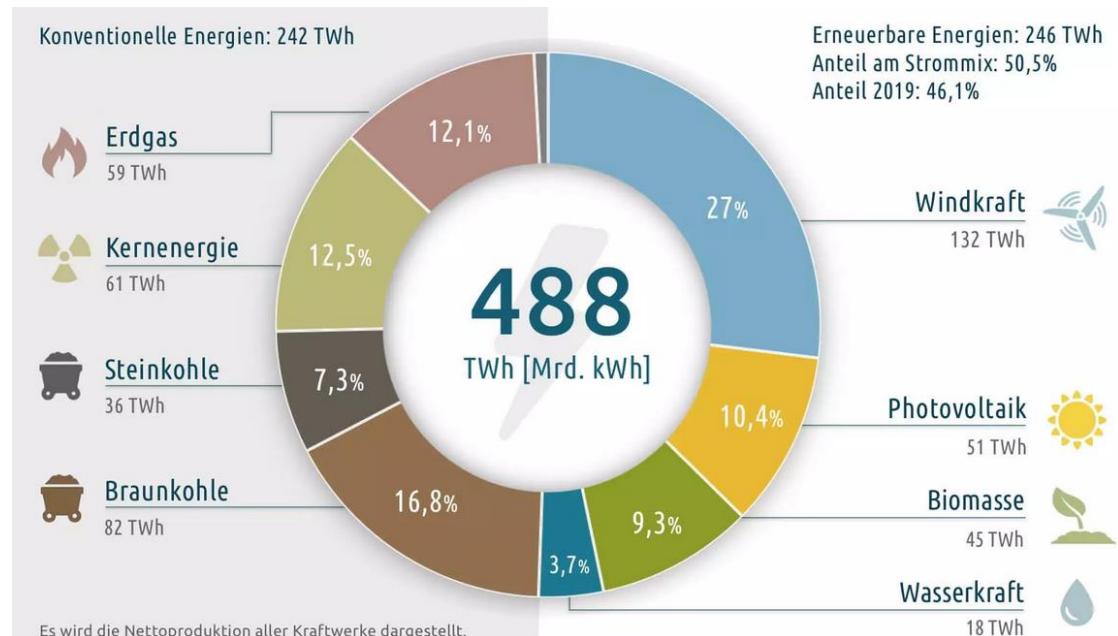
**Question 1:** how much electricity does Germany consume?

→ 488 TWh (2019)

→ 5.9 MWh per person per year

**Question 2:** how much CO<sub>2eq</sub> does this emit?

→ 400kg CO<sub>2eq</sub> / MWh



\*each country has a different mix of electricity sources

# Germany + electricity consumption

**Question 1:** how much electricity does Germany consume?

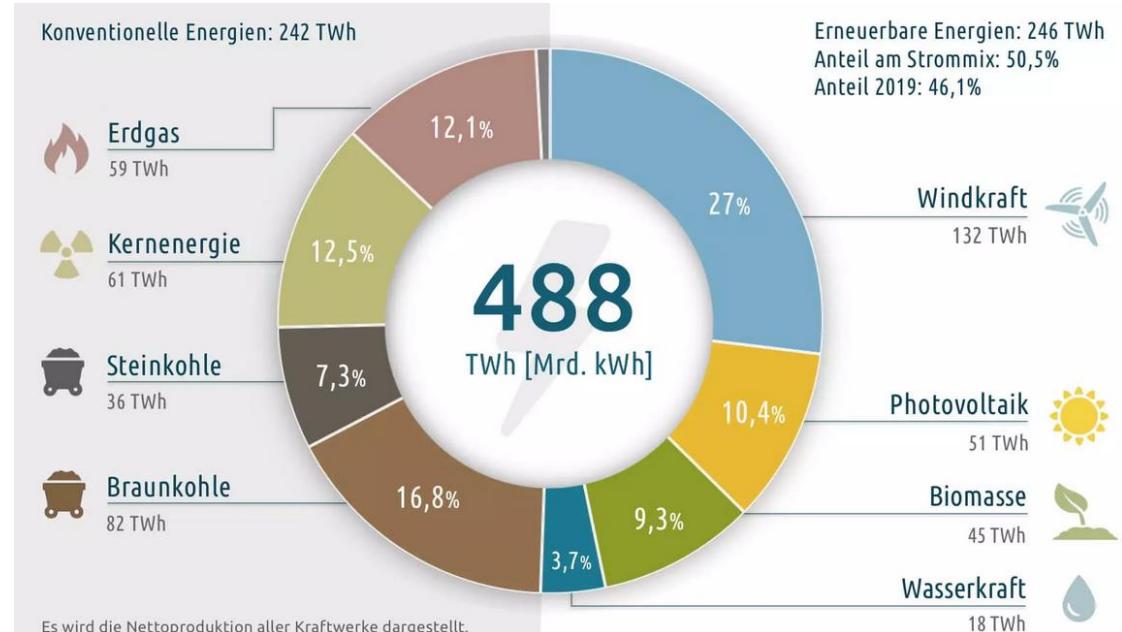
→ 488 TWh (2019)

→ 5.9 MWh per person per year

**Question 2:** how much CO<sub>2eq</sub> does this emit?

→ 400kg CO<sub>2eq</sub> / MWh

→ 2.4 tCO<sub>2eq</sub> per person per year

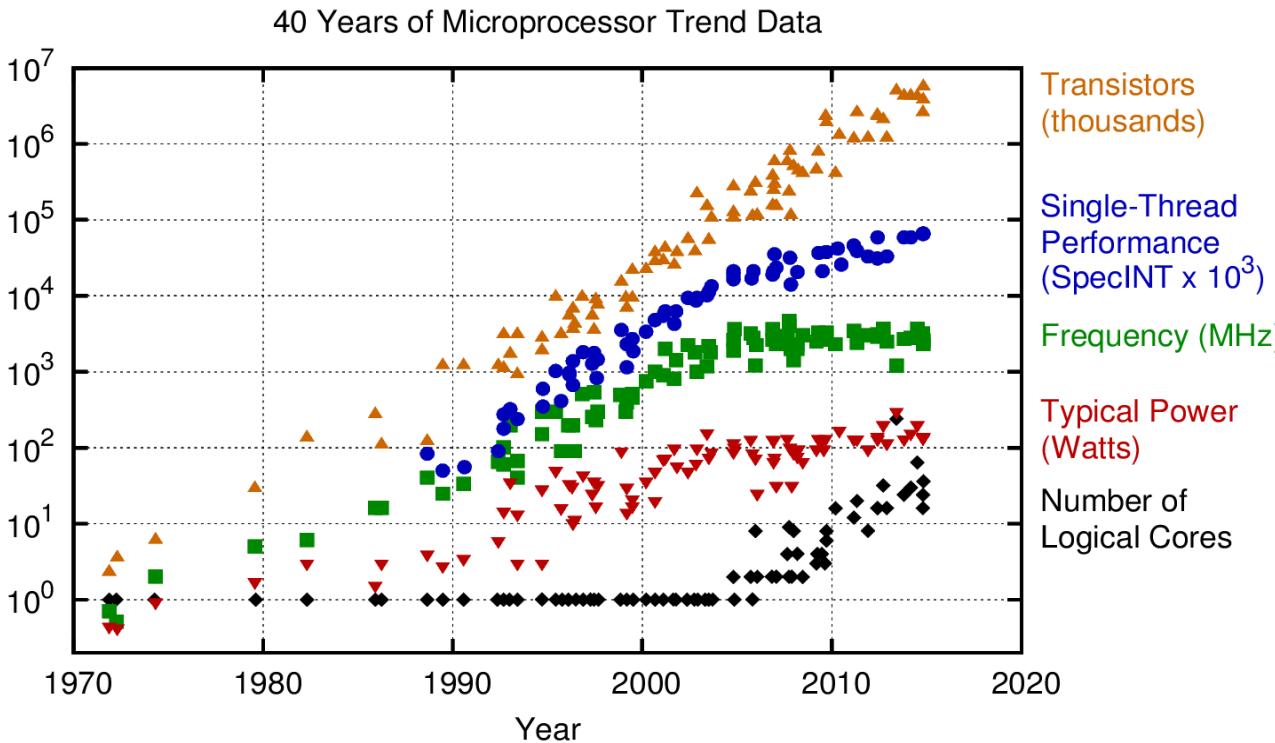


\*each country has a different mix of electricity sources

# Computer physical boundaries



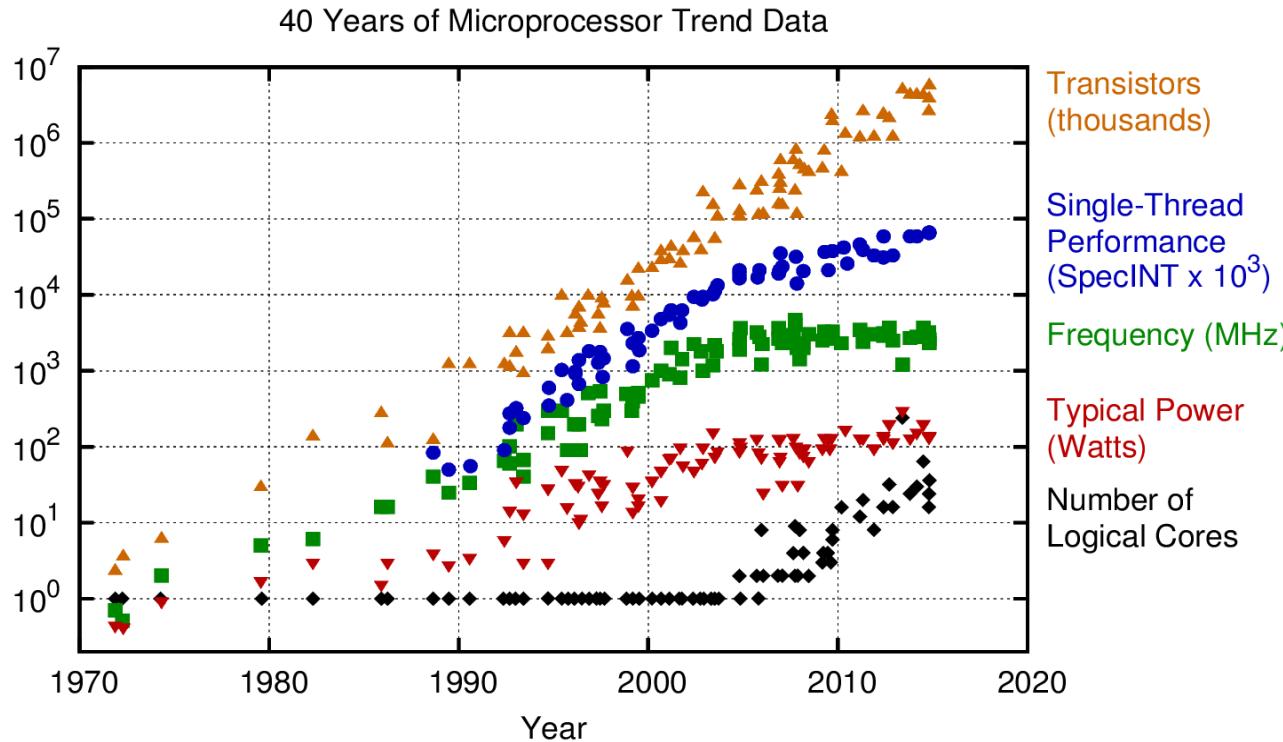
# Computer physical boundaries



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2015 by K. Rupp

# Computer physical boundaries

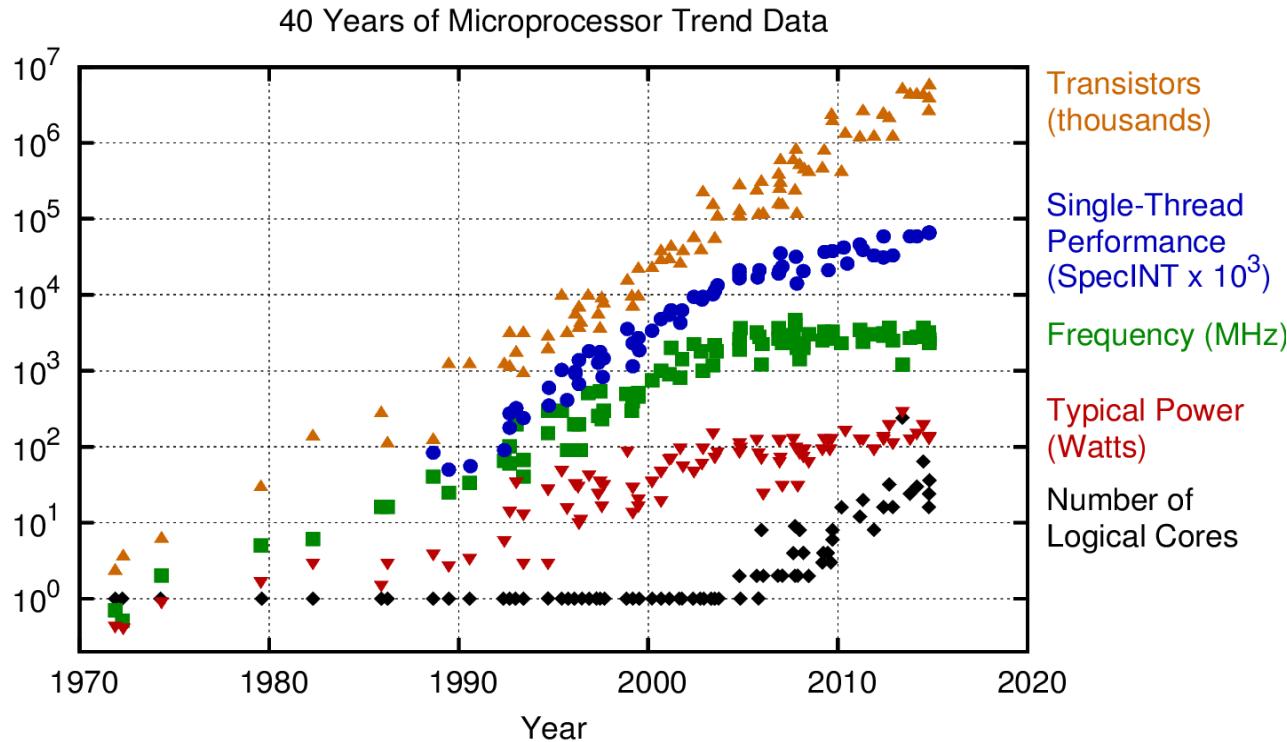
- Frequency stalled
  - Heat dissipation
  - Leak current
  - $\propto$  power increase



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2015 by K. Rupp

# Computer physical boundaries

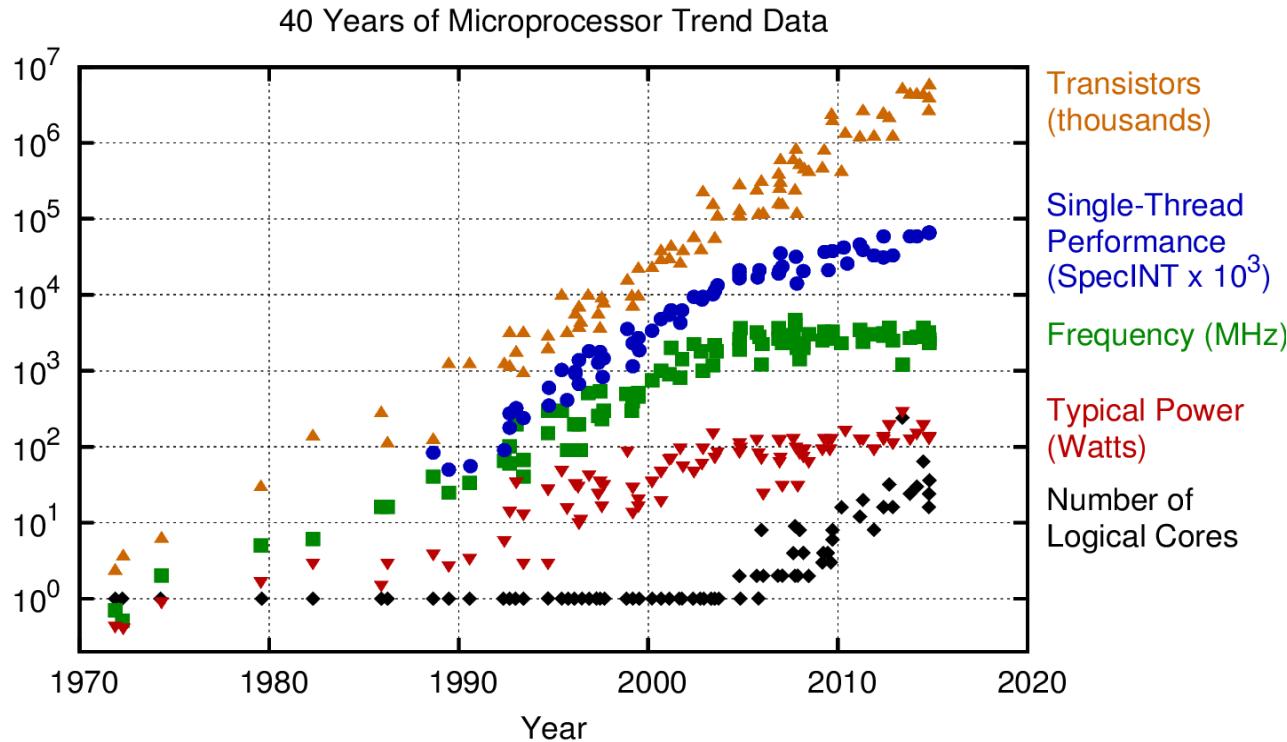
- Frequency stalled
  - Heat dissipation
  - Leak current
  - $\propto$  power increase
- Parallelism it is



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2015 by K. Rupp

# Computer physical boundaries

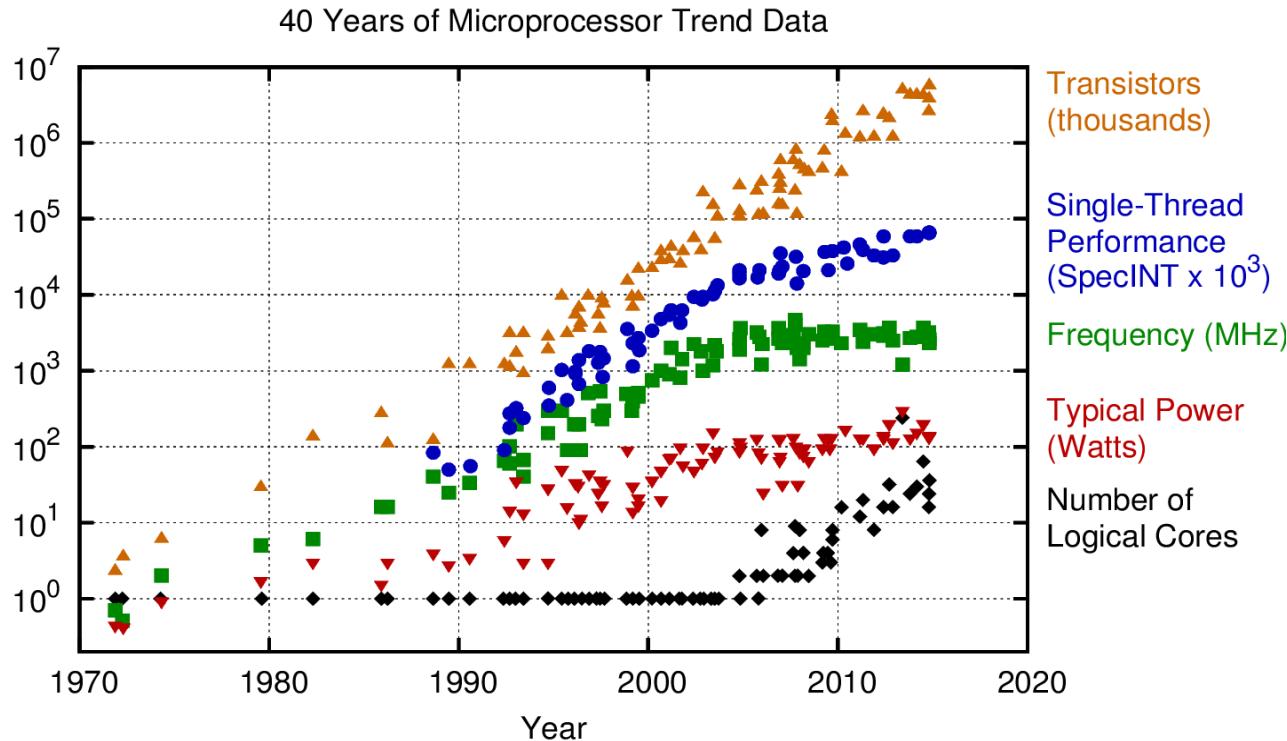
- Frequency stalled
  - Heat dissipation
  - Leak current
  - $\propto$  power increase
- Parallelism it is
- CPUs more cores



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2015 by K. Rupp

# Computer physical boundaries

- Frequency stalled
  - Heat dissipation
  - Leak current
  - $\propto$  power increase
- Parallelism it is
- CPUs more cores
- GPUs



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2015 by K. Rupp

copied from David Mytton (@davidmytton on twitter)

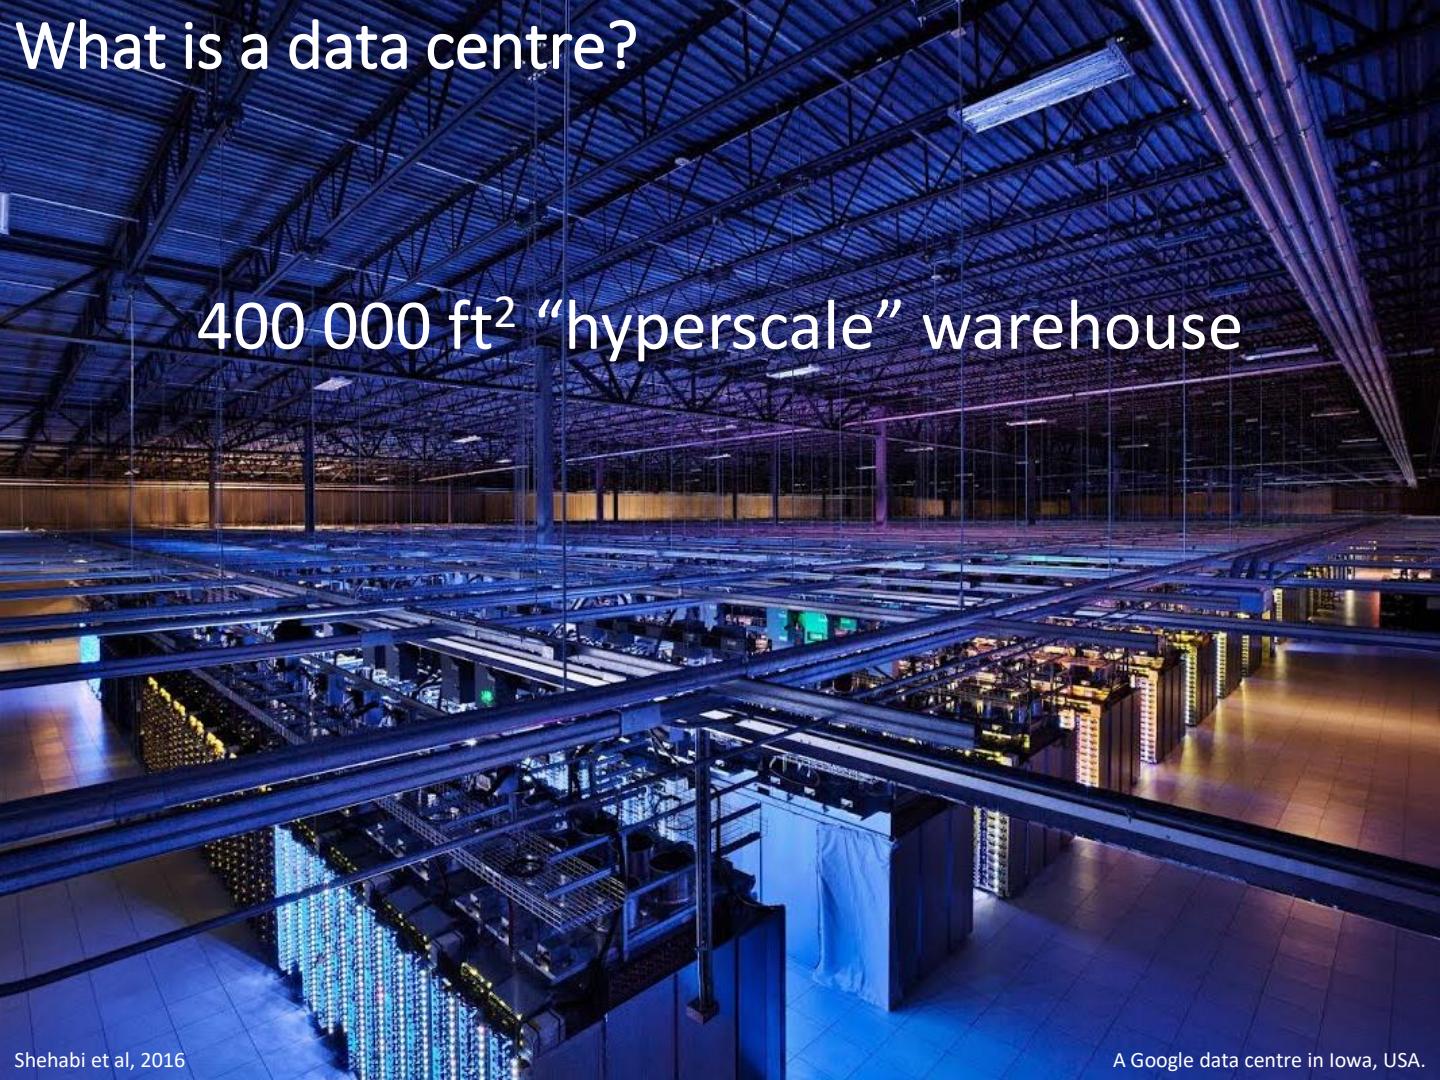
## Energy (and water) consumption in IT



A Google data centre in Hamina, Finland.

# What is a data centre?

400 000 ft<sup>2</sup> “hyperscale” warehouse



# Data centre components

1. Servers
2. Storage
3. Network
4. Infrastructure



A Google data centre in Oregon, USA.

# Component 1: Servers

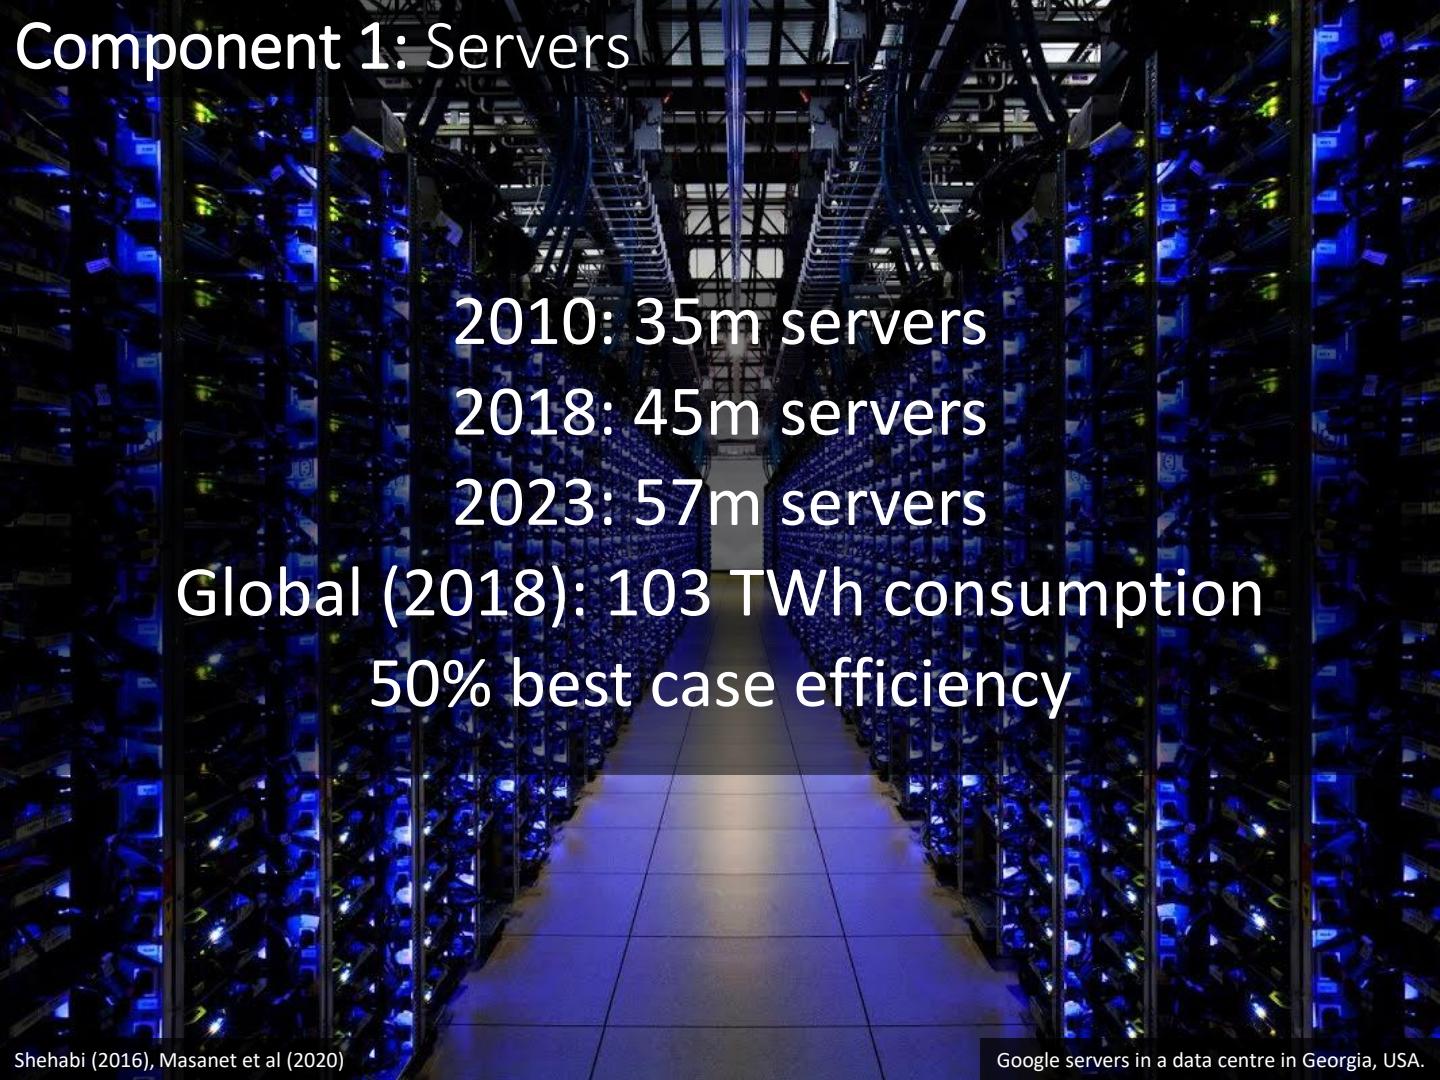
2010: 35m servers

2018: 45m servers

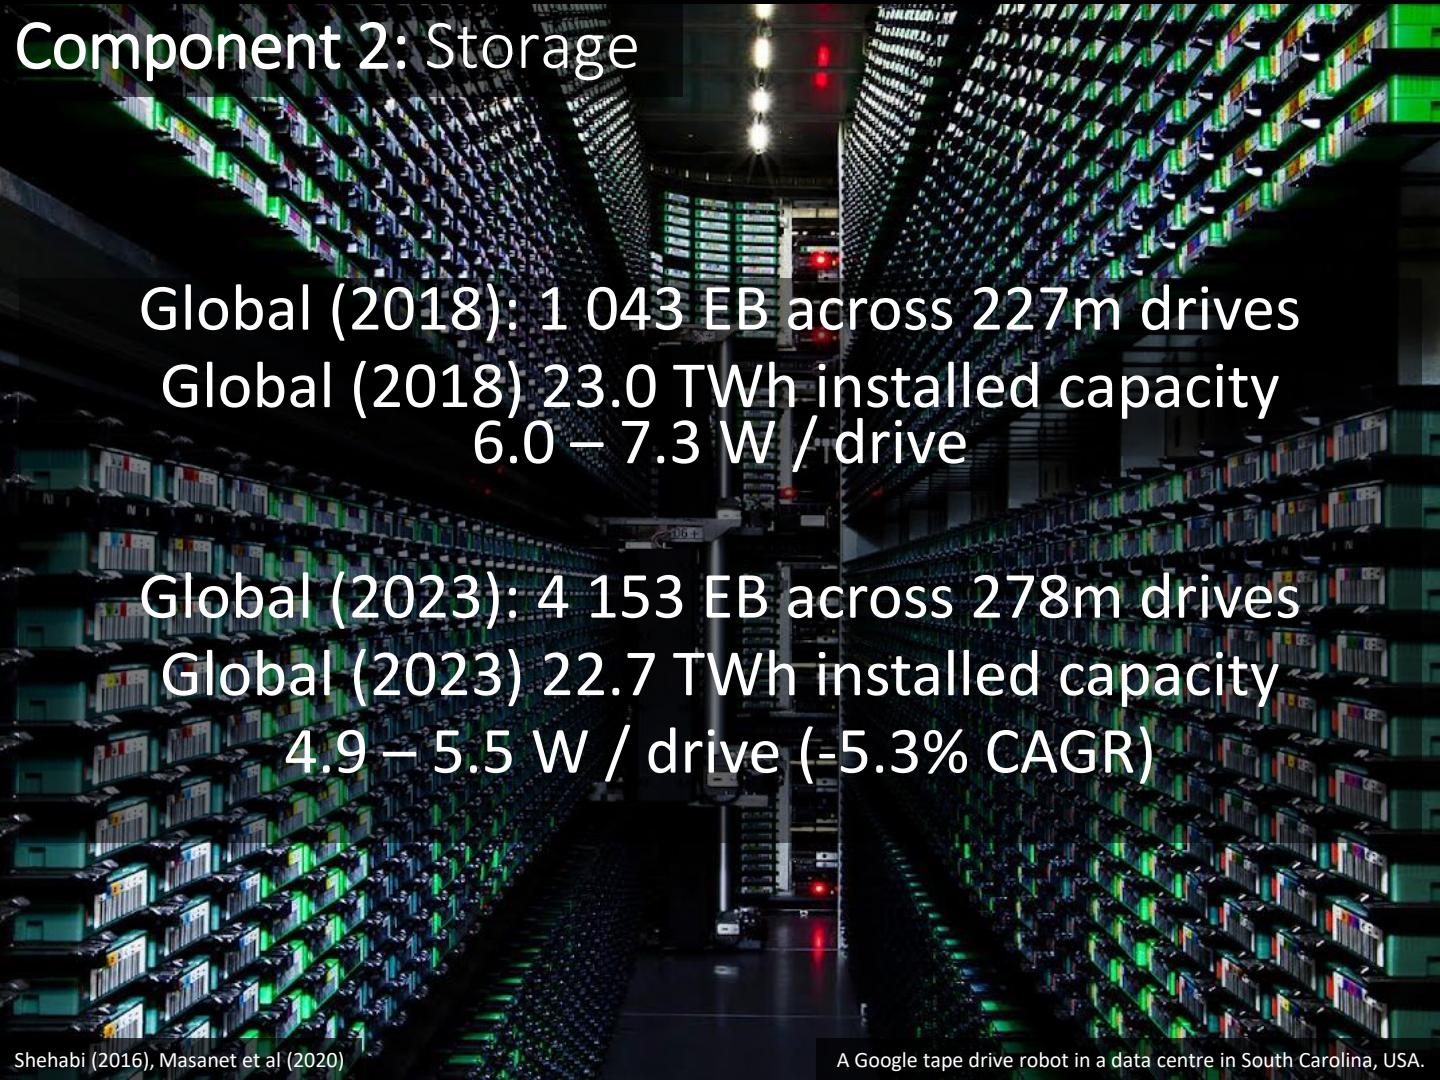
2023: 57m servers

Global (2018): 103 TWh consumption

50% best case efficiency



## Component 2: Storage



Global (2018): 1 043 EB across 227m drives

Global (2018) 23.0 TWh installed capacity  
6.0 – 7.3 W / drive

Global (2023): 4 153 EB across 278m drives

Global (2023) 22.7 TWh installed capacity  
4.9 – 5.5 W / drive (-5.3% CAGR)

# Component 3: Network

2015: 0.06 kWh/GB, decreasing 50% every 2yr

But excludes:

Mobile = 20% of internet traffic by 2022, growing 43%/yr

Internal traffic = doubling every 12-15 months

# Component 4: Infrastructure

## Power Usage Effectiveness

$$PUE = \frac{\text{Total Facility Energy}}{\text{IT Equipment Energy}} = 1 + \frac{\text{Non Facility Energy}}{\text{IT Equipment Energy}}$$



# Component 4: Infrastructure

## Power Usage Effectiveness

$$PUE = \frac{\text{Total Facility Energy}}{\text{IT Equipment Energy}} = 1 + \frac{\text{Non Facility Energy}}{\text{IT Equipment Energy}}$$

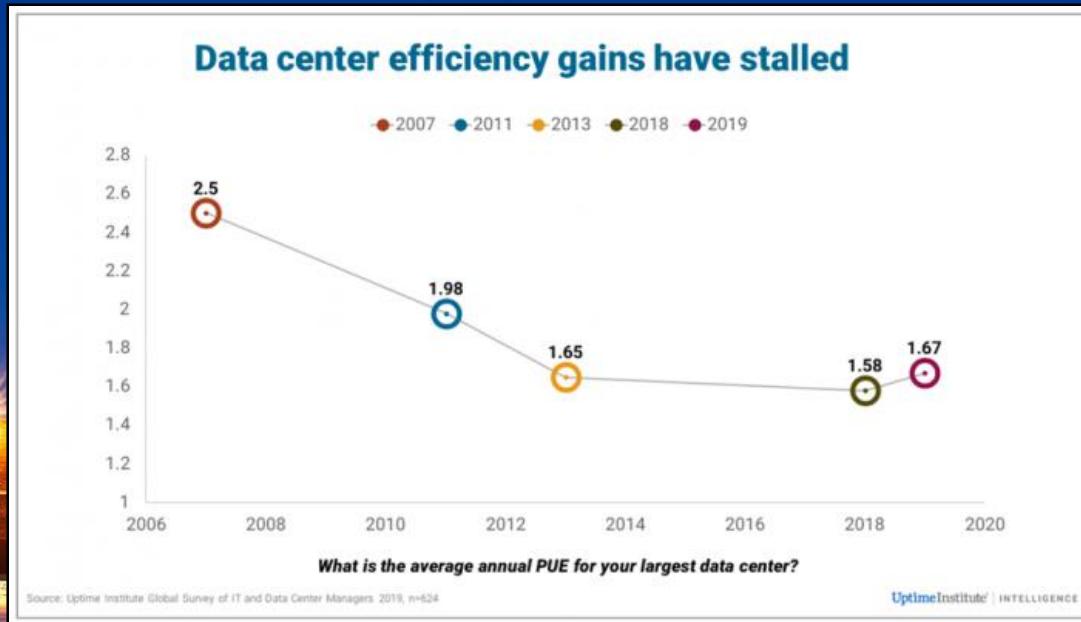
1.0 = 100% power to IT equipment

Industry average: 1.67



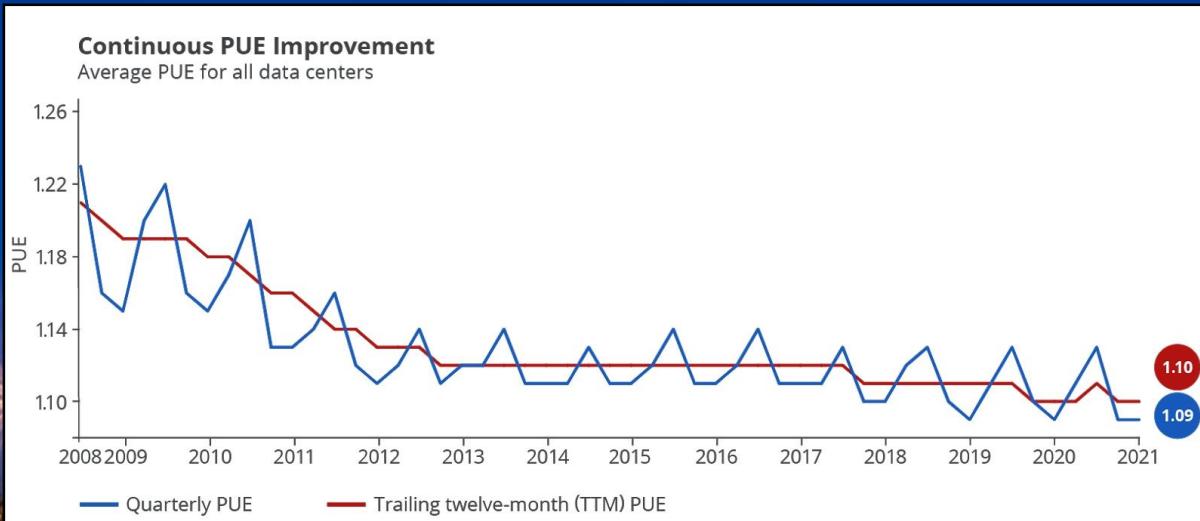
# Component 4: Infrastructure

## Power Usage Effectiveness (Industry wide)

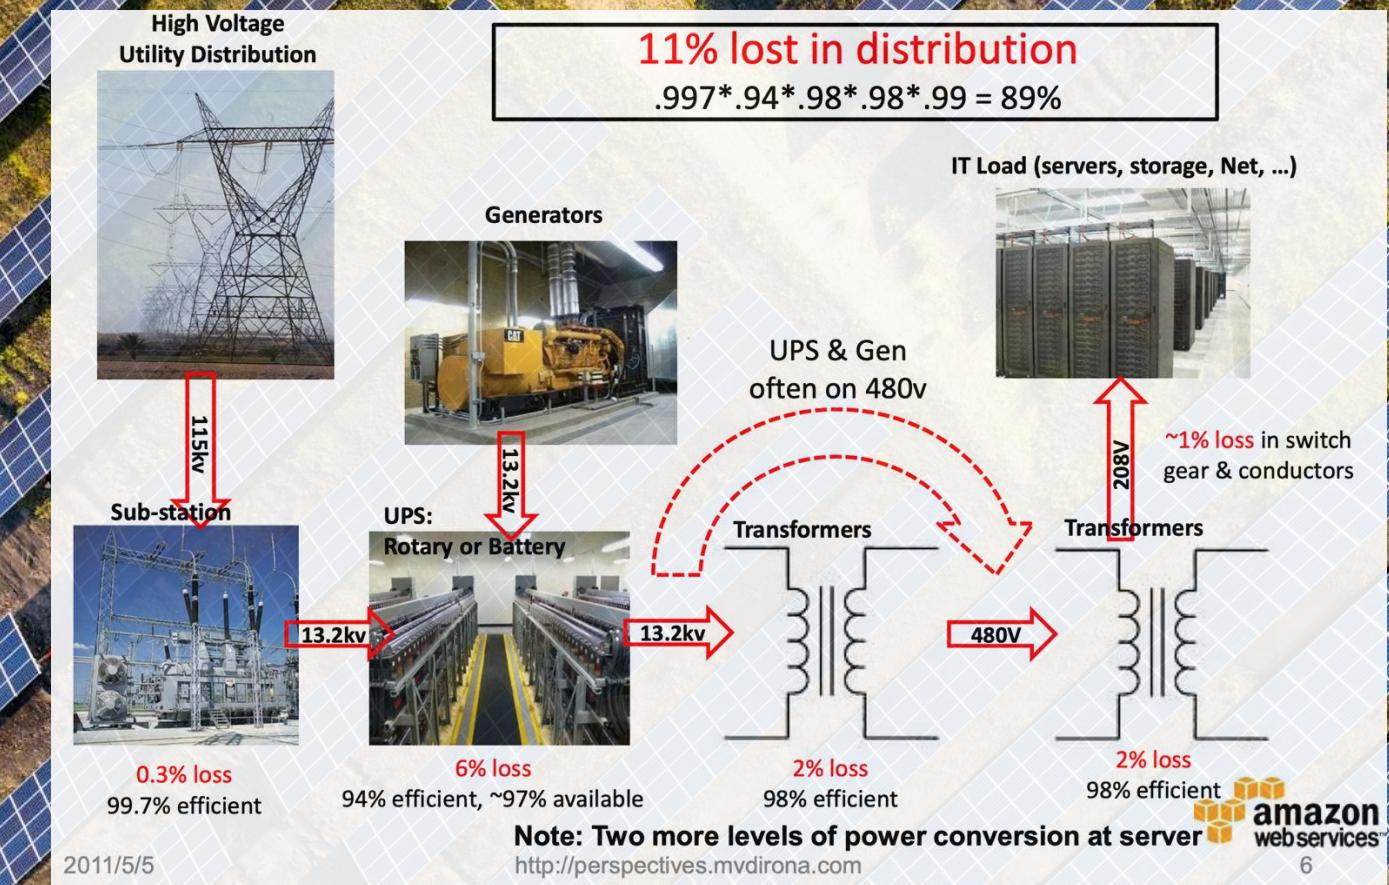


# Component 4: Infrastructure

## Power Usage Effectiveness (Google)

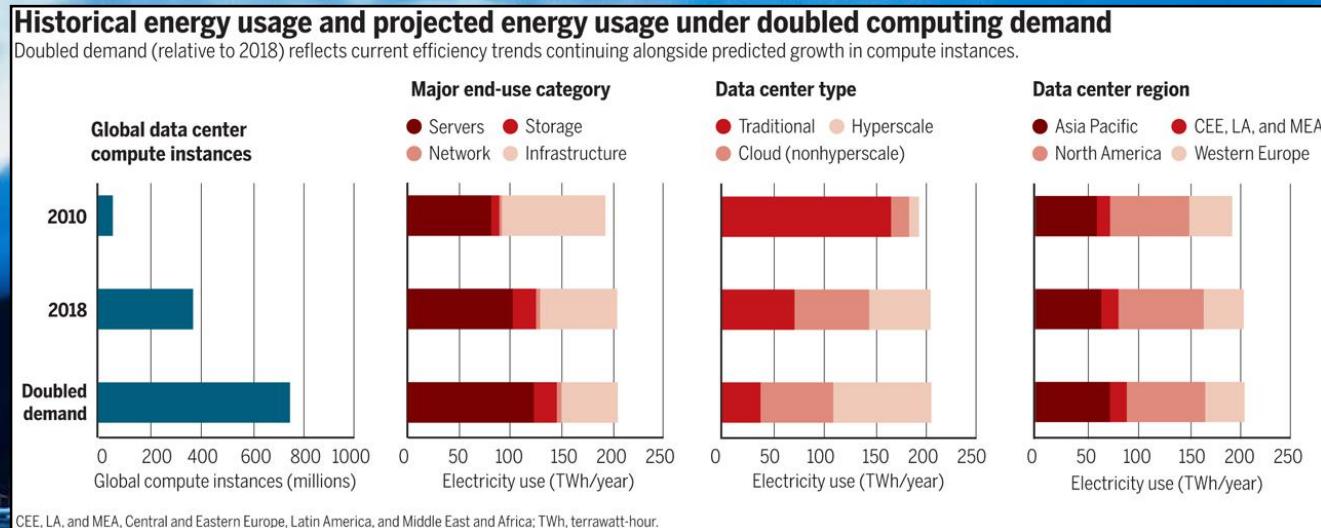


# Power distribution



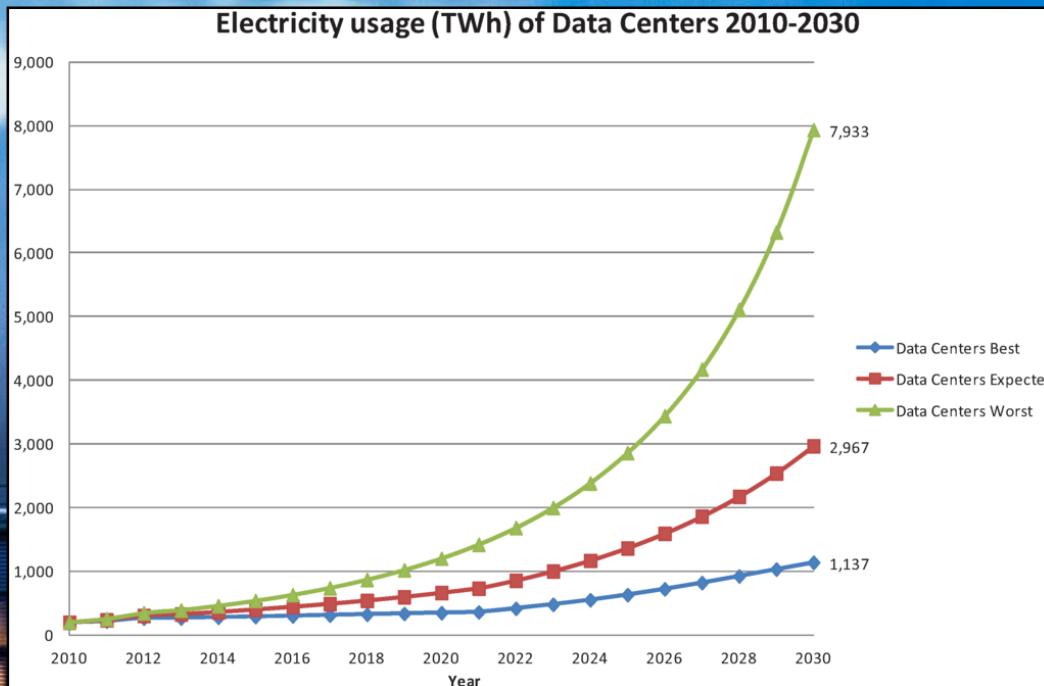
# How much energy? Position 1

2018 = 200 TWh/year & growth has plateaued

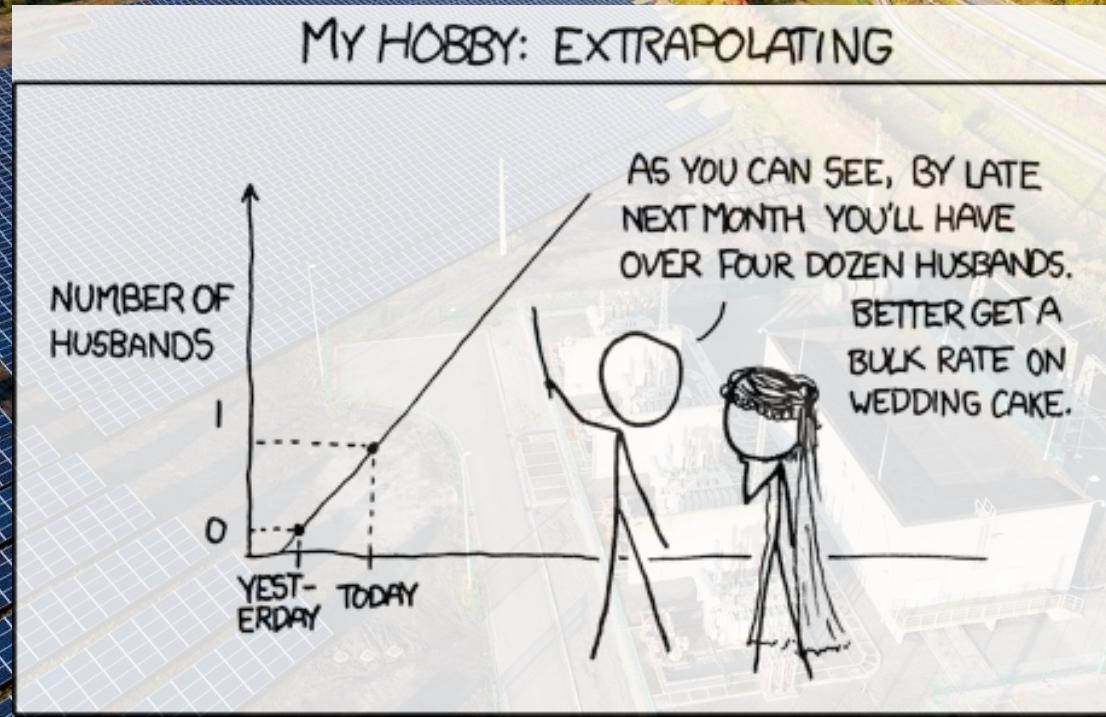
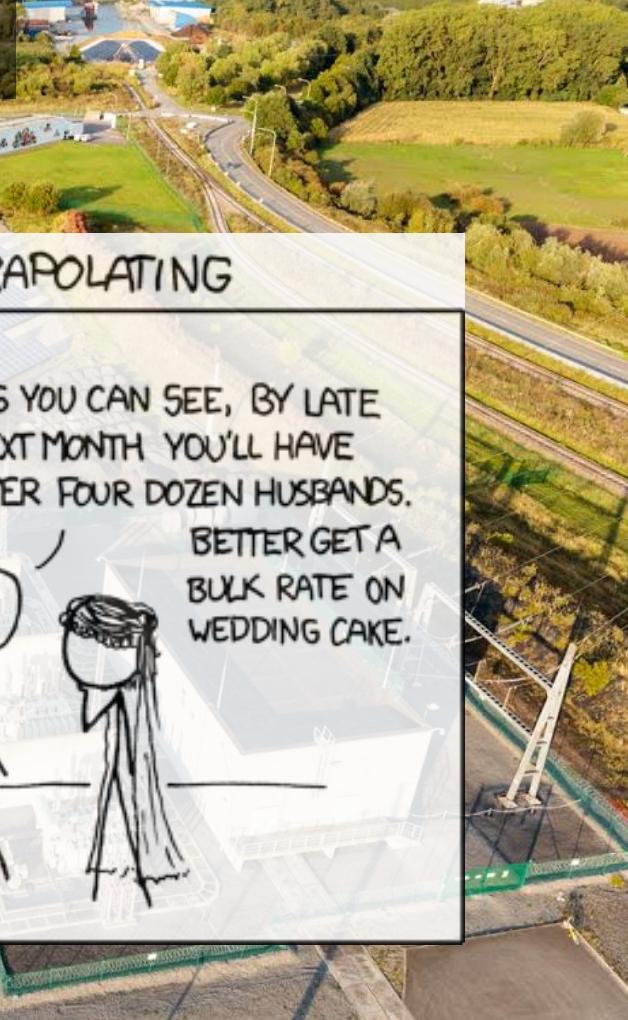


# How much energy? Position 2

2018 = 600 TWh/year & growth is massive

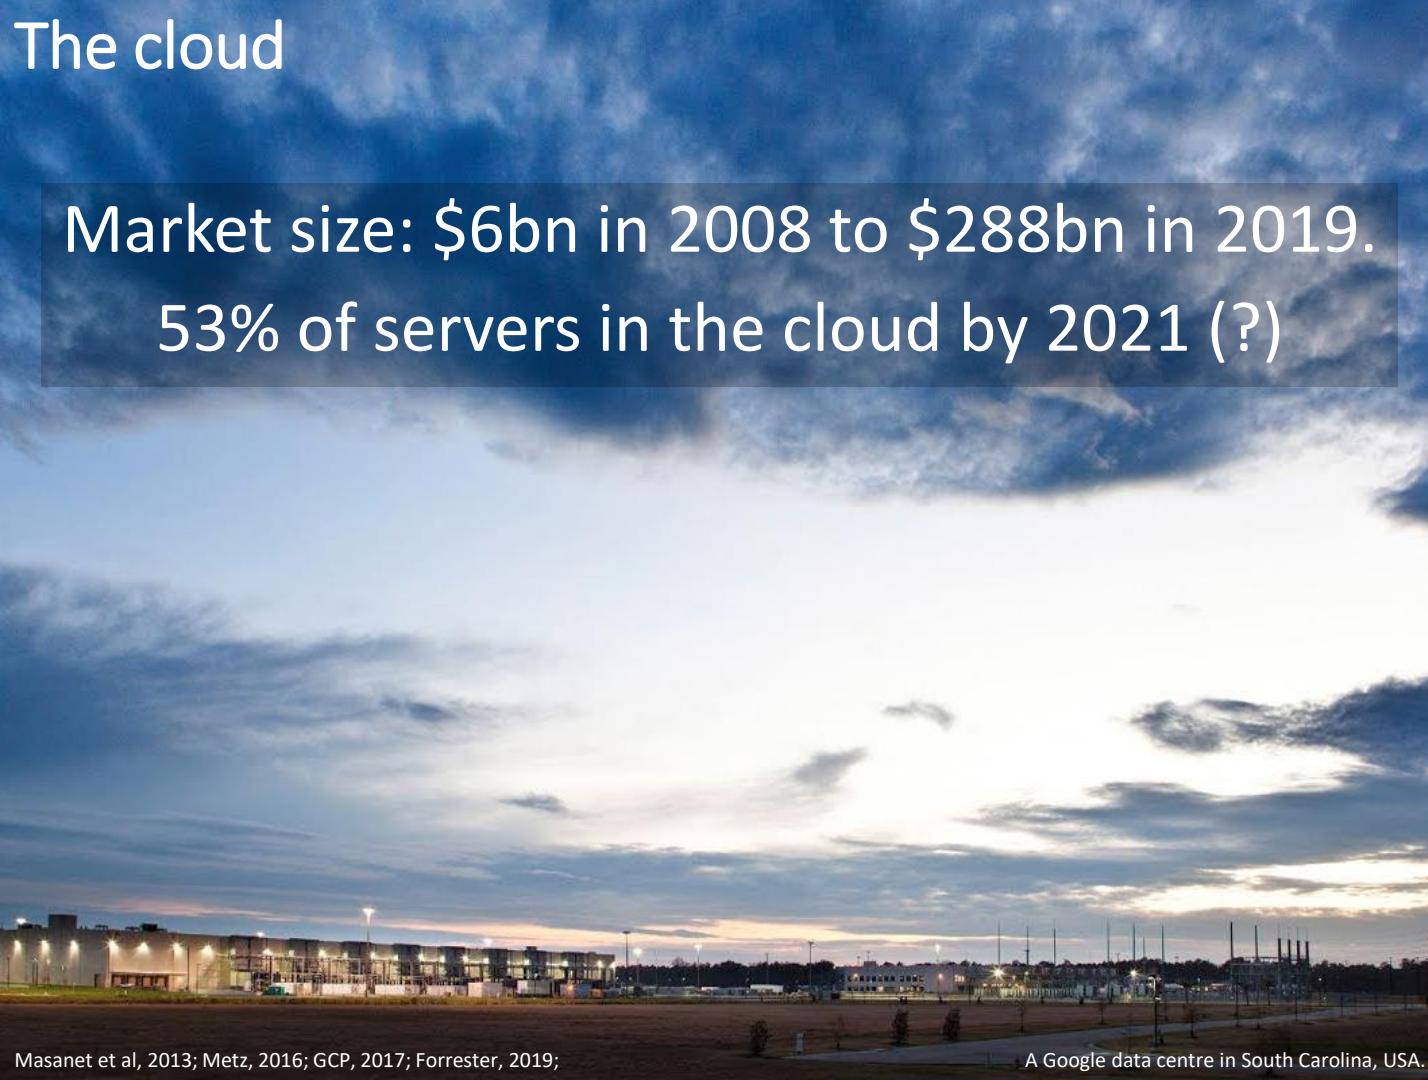


# Massive growth? 7,933 TWh?



# The cloud

Market size: \$6bn in 2008 to \$288bn in 2019.  
53% of servers in the cloud by 2021 (?)

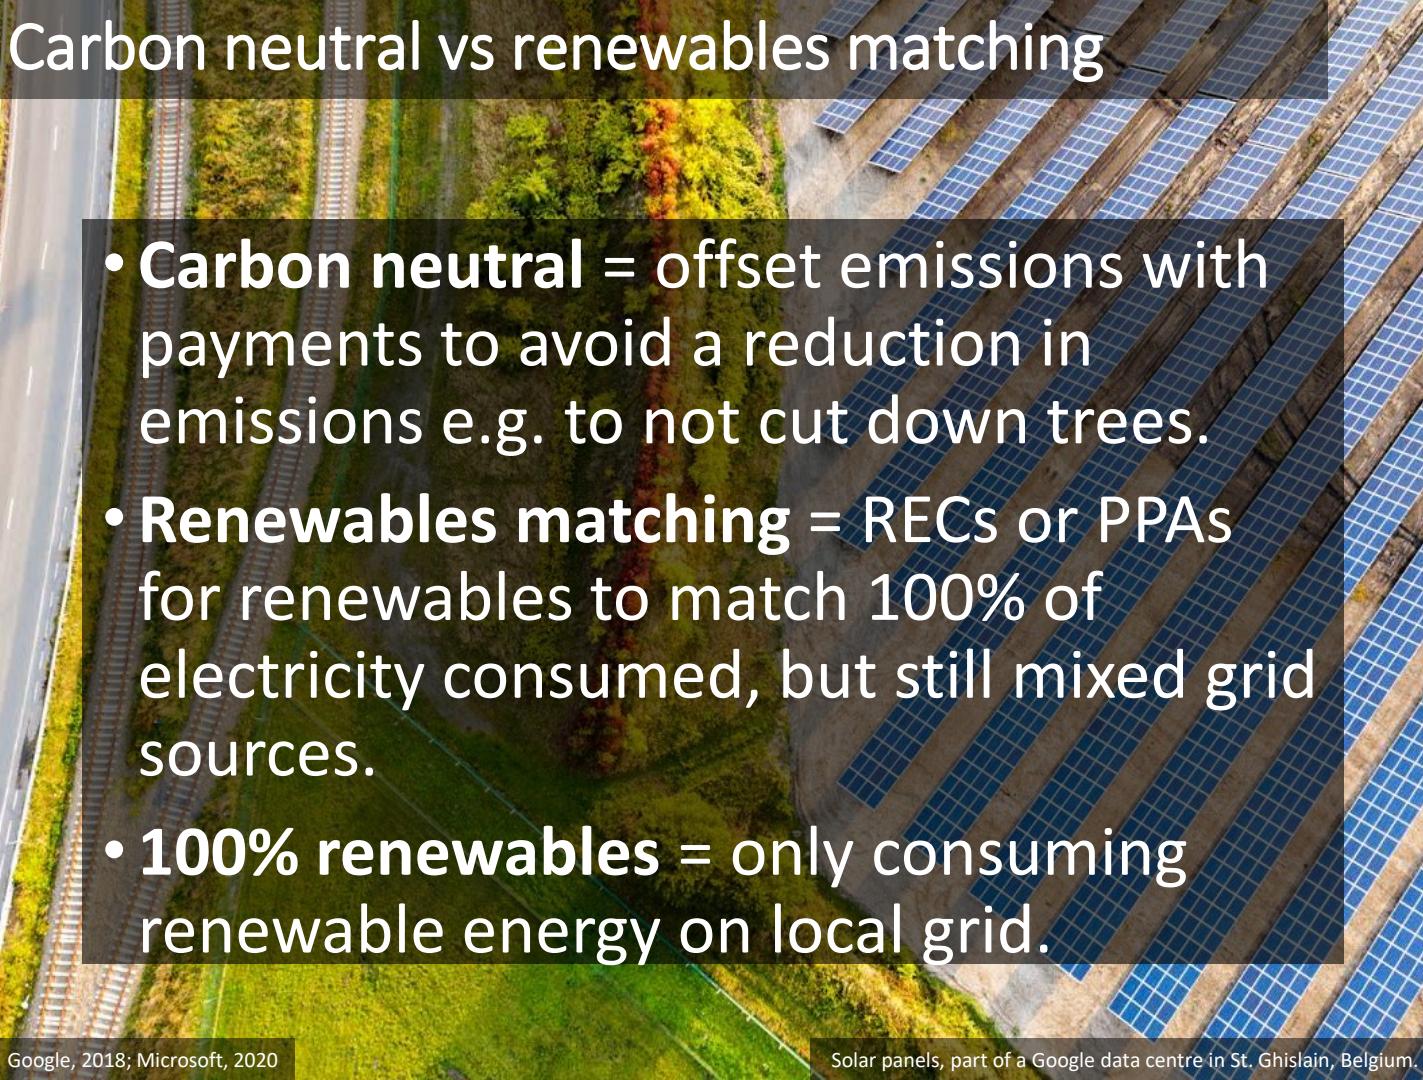


# The big 3 cloud vendors

	Amazon	Google	Microsoft
Average PUE	Under 1.2	1.10 (Q1 2021)	1.125 (only new data centres)
Power mix average carbon intensity	393 gCO <sub>2</sub> e/kWh (2015)	~61 gCO <sub>2</sub> e/kWh (2019)	?
Carbon neutral	5 of 24 regions	Globally since 2007	Globally since 2012
Renewables matching	By 2025	Since 2017	By 2025
100% Renewables	?	By 2030	Sweden in 2021 Globally by 2030

# Carbon neutral vs renewables matching

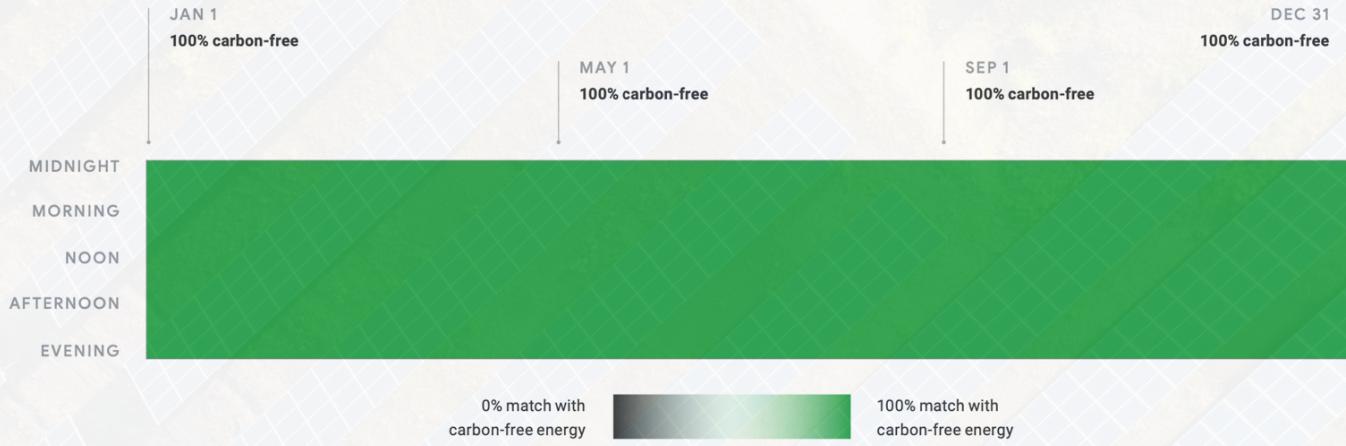
- **Carbon neutral** = offset emissions with payments to avoid a reduction in emissions e.g. to not cut down trees.
- **Renewables matching** = RECs or PPAs for renewables to match 100% of electricity consumed, but still mixed grid sources.
- **100% renewables** = only consuming renewable energy on local grid.



# Carbon neutral vs renewables matching

## Every hour of electricity use at a 24x7 carbon-free data center

A data center that is 100% matched with carbon-free energy around the clock — 24 hours a day, 7 days a week, 365 days per year



# Carbon neutral vs renewables matching

## Every hour of electricity use at Finland data center

Google's purchases of wind power in the Nordic region, combined with significant nuclear and hydropower on the grid, mean that in most hours this data center is 100% matched with carbon-free energy

Overall in 2017, 97% of this data center's electricity use was matched on an hourly basis with carbon-free sources.

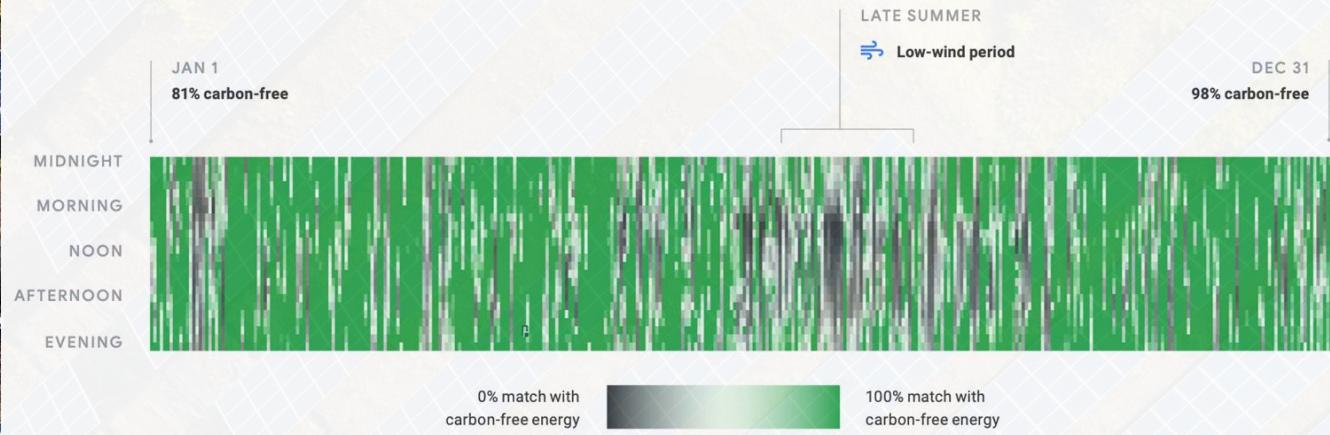


# Carbon neutral vs renewables matching

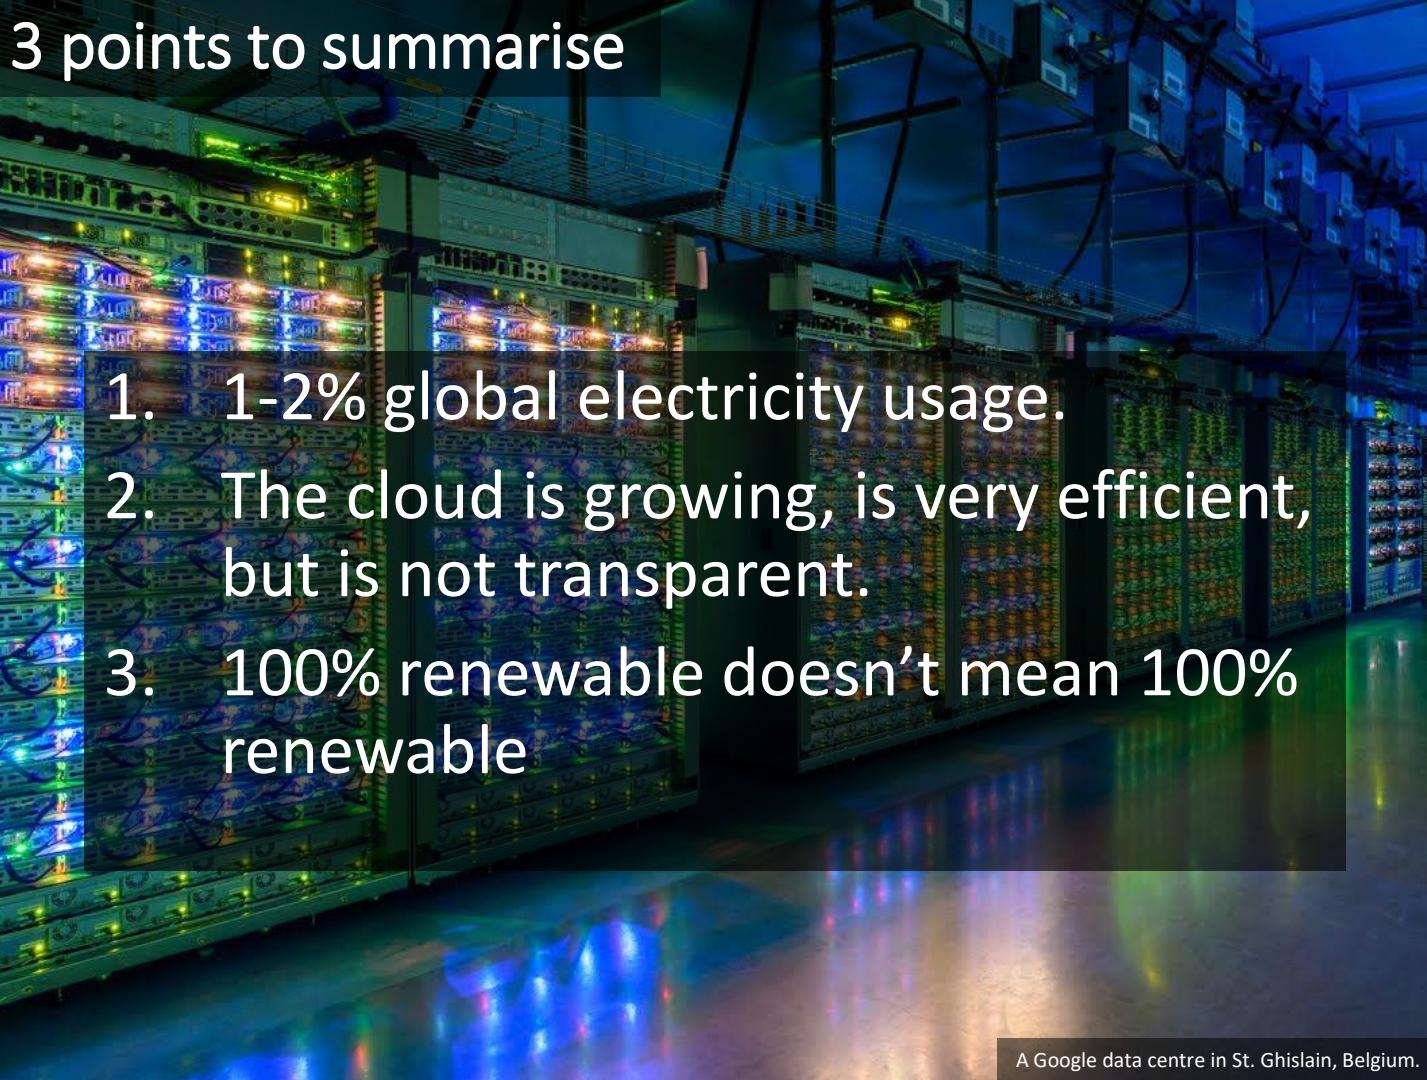
## Every hour of electricity use at Iowa data center

Although our Iowa data center achieved 100% carbon-free energy during the majority of hours in 2017, there is also a recurring reliance on carbon-based power – most notably in late summer, when wind speeds decline.

Overall in 2017, 74% of this data center's electricity use was matched on an hourly basis with carbon-free sources.



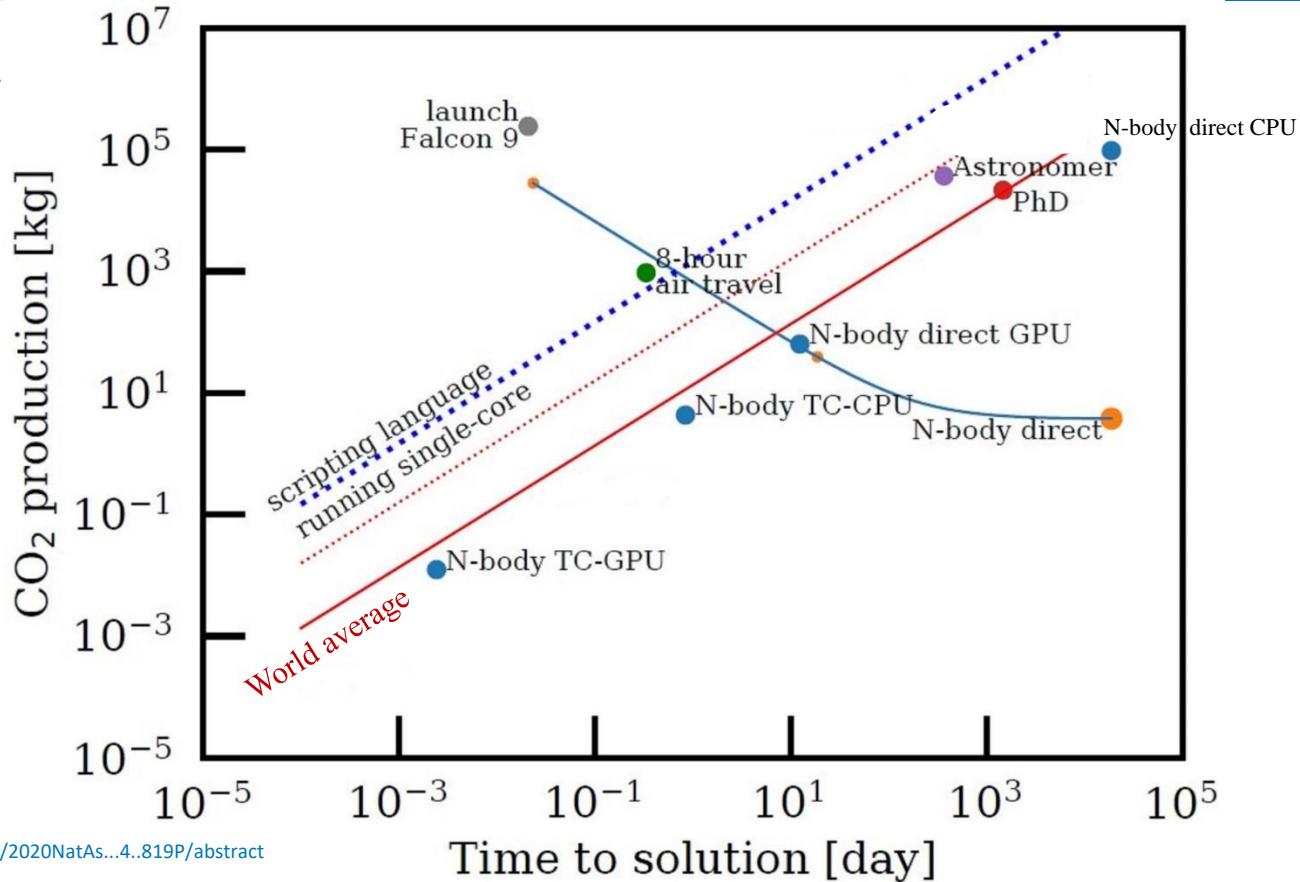
# 3 points to summarise

- 
- 1. 1-2% global electricity usage.
  - 2. The cloud is growing, is very efficient, but is not transparent.
  - 3. 100% renewable doesn't mean 100% renewable

# Increasing computing efficiency



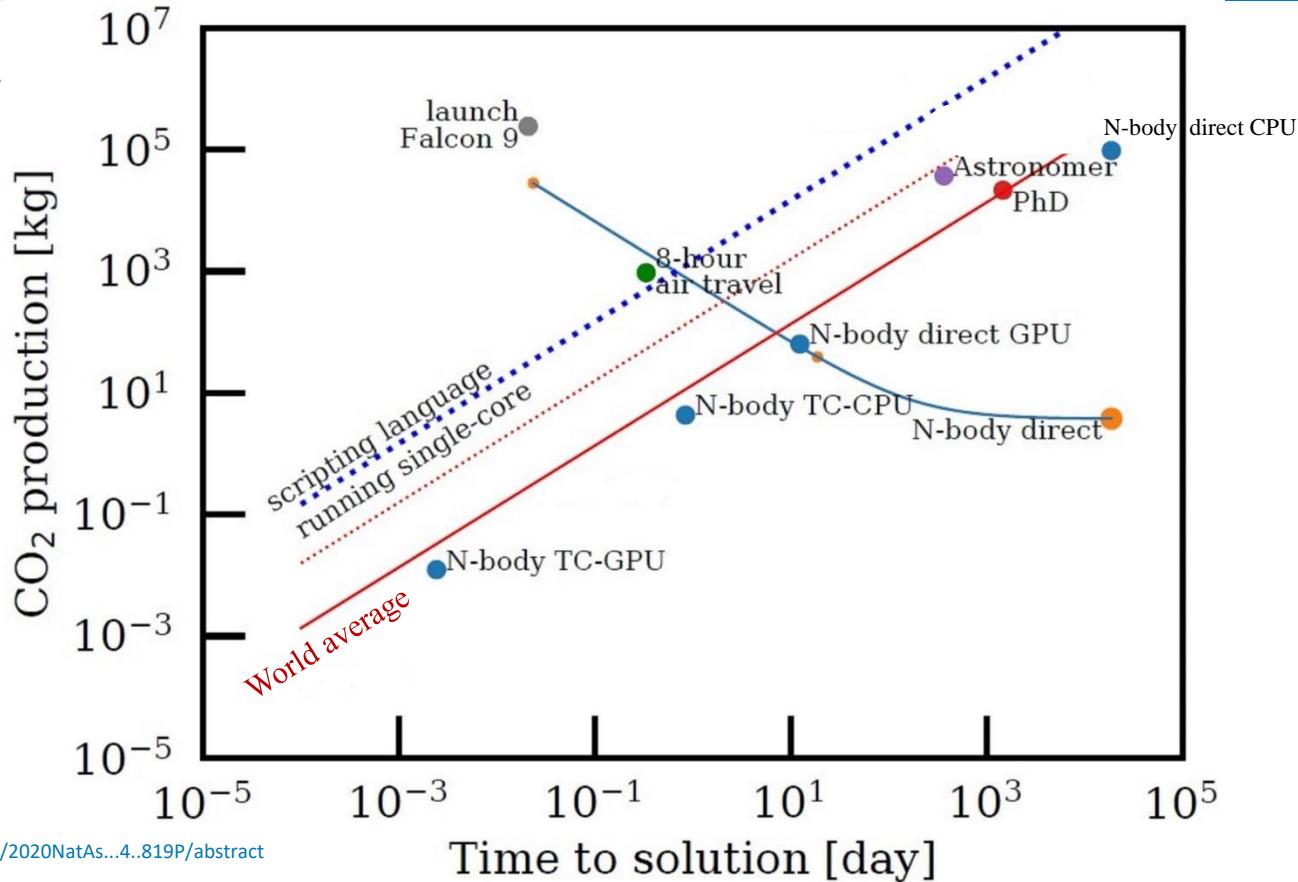
## ■ Calculate wisely



# Increasing computing efficiency



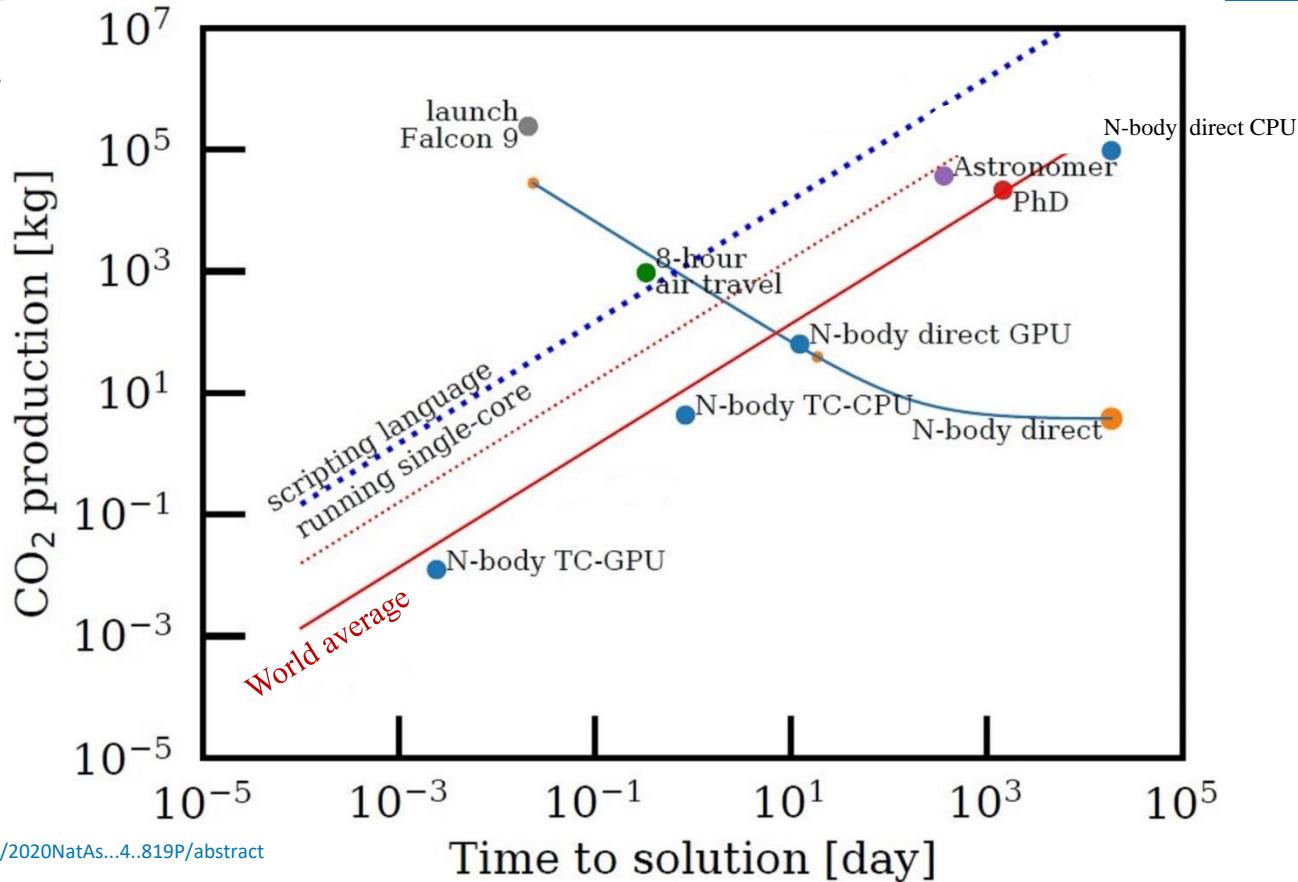
- Calculate wisely
- hardware



# Increasing computing efficiency



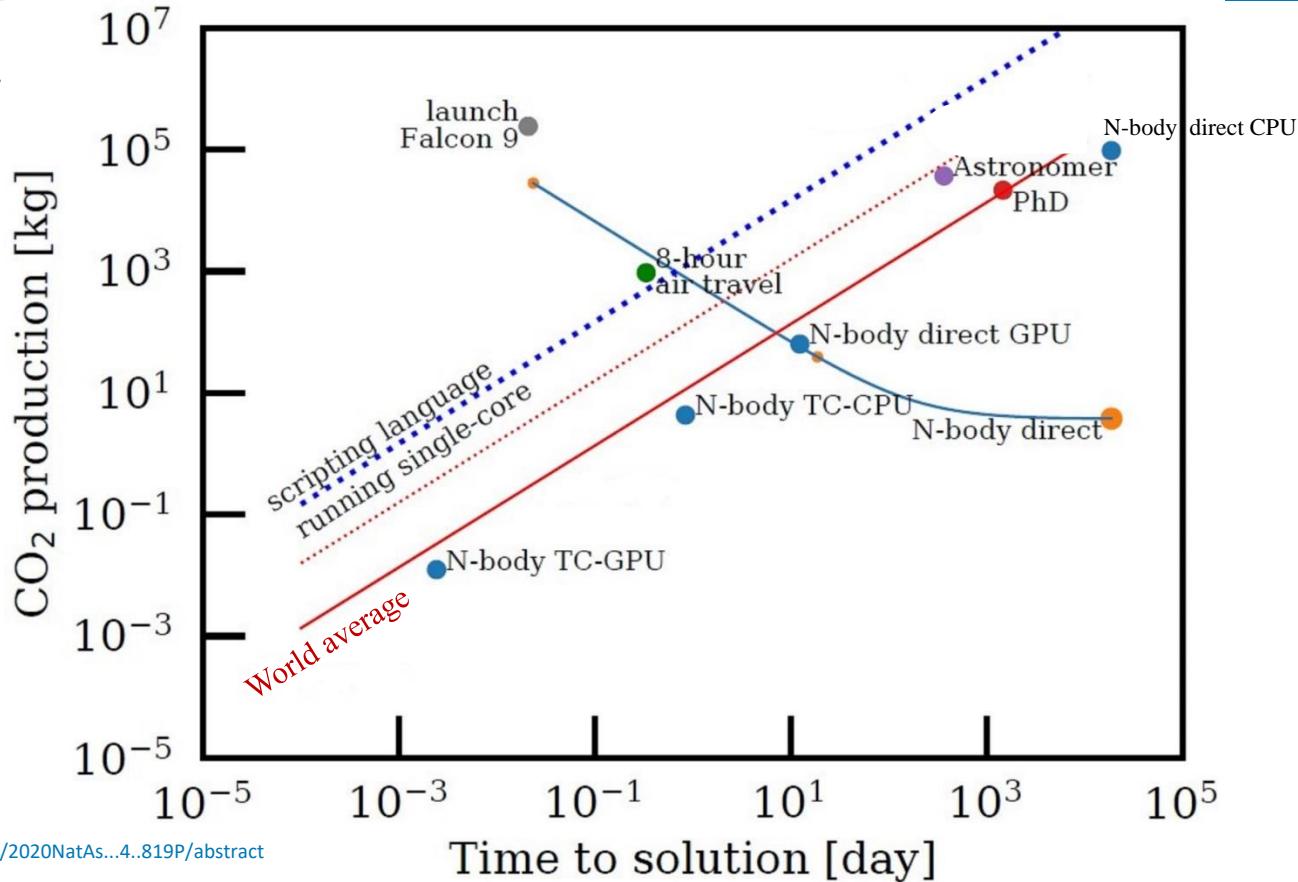
- Calculate wisely
- hardware
- code



# Increasing computing efficiency



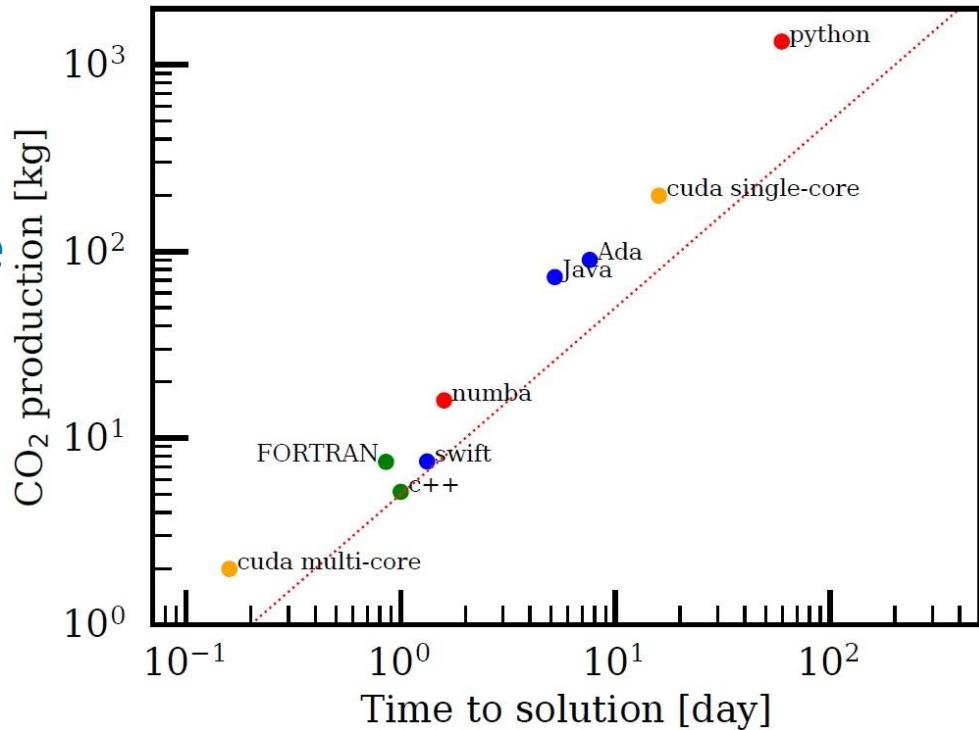
- Calculate wisely
- hardware
- code
- parallelism



# Increasing computing efficiency



- Python vs. Other languages
- Your own time
  - I had to use JAVA für Gaia
  - With python I would have been able to optimize my code further
- Optimize when using heavily



# Increasing computing efficiency



- Improve your code

- Improve your code
  - <https://github.com/uschpc/workshop-hpc-python>
  - Profiling
  - Parallel IO

- Improve your code
  - <https://github.com/uschpc/workshop-hpc-python>
  - Profiling
  - Parallel IO
- Write important bits in fast languages

- Improve your code
  - <https://github.com/uschpc/workshop-hpc-python>
  - Profiling
  - Parallel IO
- Write important bits in fast languages
- Or make your code portable between hardware
  - <https://github.com/kokkos/kokkos>

# Increasing computing efficiency



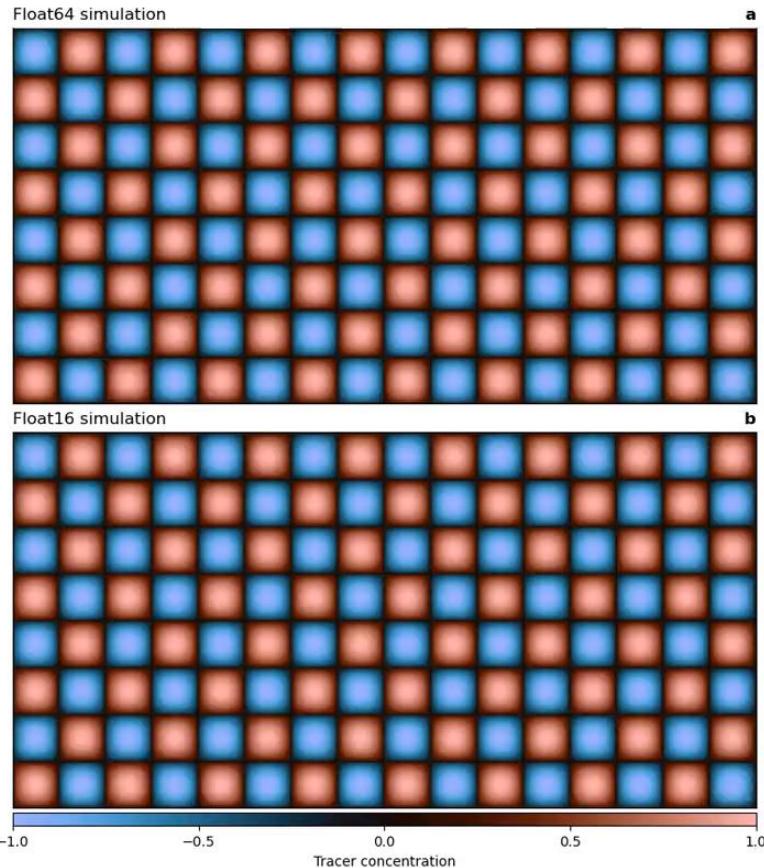
- Lower resolution

# Increasing computing efficiency



- Lower resolution
- lower accuracy
  - Klöwer+ 2022

(<https://agupubs.onlinelibrary.wiley.com/doi/10.1029/2021MS002684>)



# Increasing computing efficiency



- Storing and sending data costs energy

# Increasing computing efficiency



- Storing and sending data costs energy
- Use compression (lossy compression)

# Increasing computing efficiency



- Storing and sending data costs energy
- Use compression (lossy compression)
- Only store and send what is needed

# Increasing computing efficiency



- Storing and sending data costs energy
- Use compression (lossy compression)
- Only store and send what is needed
- Sending is approximately 10kWh per TB

# Increasing computing efficiency



- Storing and sending data costs energy
- Use compression (lossy compression)
- Only store and send what is needed
- Sending is approximately 10kWh per TB
  - 4kg of CO<sub>2</sub>
  - 3€ of electricity cost

# Increasing computing efficiency



- Work collectively and reproducibly

# Increasing computing efficiency



- Work collectively and reproducibly
  - Share infrastructure and knowledge bases

- Work collectively and reproducibly
  - Share infrastructure and knowledge bases
  - Share results, make them reproducible

- Work collectively and reproducibly
  - Share infrastructure and knowledge bases
  - Share results, make them reproducible
  - Share code and improve codes together

- Work collectively and reproducibly
  - Share infrastructure and knowledge bases
  - Share results, make them reproducible
  - Share code and improve codes together
  - Plan code development (employ and teach specialists)

- Work collectively and reproducibly
  - Share infrastructure and knowledge bases
  - Share results, make them reproducible
  - Share code and improve codes together
  - Plan code development (employ and teach specialists)
  - Make your C++ code available to python ecosystem
    - <https://docs.mpcdf.mpg.de/bnb/208.html#python-bindings-for-c-using-pybind11-and-scikit-build>

# Increasing computing efficiency



- Simplify where possible

# Increasing computing efficiency



- Simplify where possible
  - Chempy flexible one-zone chemical evolution code

Rybicki+ 2016 <https://ui.adsabs.harvard.edu/abs/2017A&26A...605A..59R/abstract>

# Increasing computing efficiency



- Simplify where possible
  - Chempy flexible one-zone chemical evolution code
  - Exploring parameter space via MCMC (1s per call)

Rybicki+ 2016 <https://ui.adsabs.harvard.edu/abs/2017A&lt;sup&gt;26A...605A..59R/abstract>

- Simplify where possible
  - Chempy flexible one-zone chemical evolution code
  - Exploring parameter space via MCMC (1s per call)
  - Oliver Philcox made a neural net interpolator of Chempy

Rybicki+ 2016 <https://ui.adsabs.harvard.edu/abs/2017A&26A...605A..59R/abstract>

Philcox + 2018 <https://ui.adsabs.harvard.edu/abs/2018ApJ...861...40P/abstract>

- Simplify where possible
  - Chempy flexible one-zone chemical evolution code
  - Exploring parameter space via MCMC (1s per call)
  - Oliver Philcox made a neural net interpolator of Chempy  
→ 1ms per call and differentiation for free

Rybicki+ 2016 <https://ui.adsabs.harvard.edu/abs/2017A&26A...605A..59R/abstract>

Philcox + 2018 <https://ui.adsabs.harvard.edu/abs/2018ApJ...861...40P/abstract>

# Increasing computing efficiency



- Simplify where possible
  - Chempy flexible one-zone chemical evolution code
  - Exploring parameter space via MCMC (1s per call)
  - Oliver Philcox made a neural net interpolator of Chempy
    - 1ms per call and differentiation for free
    - Gradient descent method exploring hundreds of one-zones with exclusive and shared variables marginalizing over uncertainties in the model

Rybicki+ 2016 <https://ui.adsabs.harvard.edu/abs/2017A&#26A...605A..59R/abstract>

Philcox + 2018 <https://ui.adsabs.harvard.edu/abs/2018ApJ...861...40P/abstract>

Philcox & Rybicki 2019 <https://ui.adsabs.harvard.edu/abs/2019ApJ...887....9P/abstract>

- Replace expensive computations where possible
  - By interpolators
  - By approximate models (if accuracy still sufficient for the task)

- Raise awareness

- Raise awareness
  - E.g. acknowledge the CO<sub>2</sub> footprint of your calculations

butions of all these computing allocations. All simulations performed for this study including test runs and failed simulations used about 2.5M core-hours which under the assumption of 30W per core-hour and a CO<sub>2</sub> intensity of electricity of  $\sim 600 \text{ kgCO}_2/\text{MWh}$  amounts to a carbon footprint of  $\sim 45\text{t}$  of CO<sub>2</sub>.

Buck+ 2021

- Raise awareness
  - Acknowledge the CO2 footprint of your calculations
  - Measure/estimate your CO2 footprint
    - E.g. <https://mlco2.github.io/impact/>

- Plan ahead in projects

# Solution ideas

- Plan ahead in projects
- Example GRAND collaboration

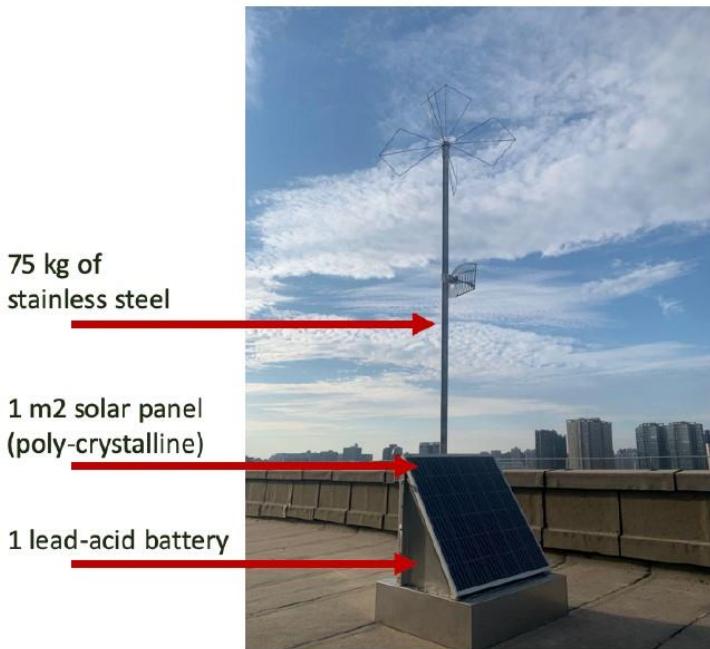


Figure 9: GRANDProto300 prototype antenna.

# Solution ideas

- Plan ahead in projects
- Example GRAND collaboration
  - 300, 10k, 200k antennas



Figure 9: GRANDProto300 prototype antenna.

# Solution ideas



## ■ Plan ahead in projects

2021

2025

203X

### GRANDProto300

### GRAND10k

### GRAND200k



Set-up

300 radio antennas  
over 200 km<sup>2</sup>  
in Western China

10,000 radio antennas  
over 10,000 km<sup>2</sup> in China

200,000 radio antennas  
over 200,000 km<sup>2</sup>  
~20 subarrays of 10,000  
antennas



People

~70 members  
(~30% strong involvement)

~100 members  
(~30% strong involvement)

~400 - 1000 members



Emission  
[tCO<sub>2</sub>e/yr]

Digital: 338 (69%)  
Travel: 130 (27%)  
Hardware: 18 (4%)  
**Total: 486 tCO<sub>2</sub>e/yr**

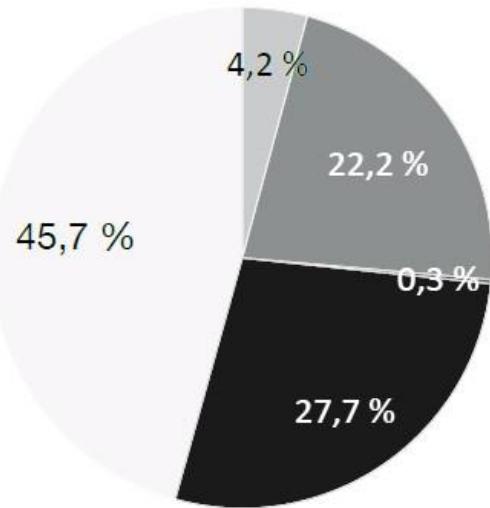
Digital: 422 (39%)  
Travel: 328 (31%)  
Hardware: 322 (30%)  
**Total: 1072 tCO<sub>2</sub>e/yr**

Hardware: 6435 (48%)  
Digital: 6007 (45%)  
Travel: 965 (7%)  
**Total: 13407 tCO<sub>2</sub>e/yr**

## ■ Plan ahead in projects

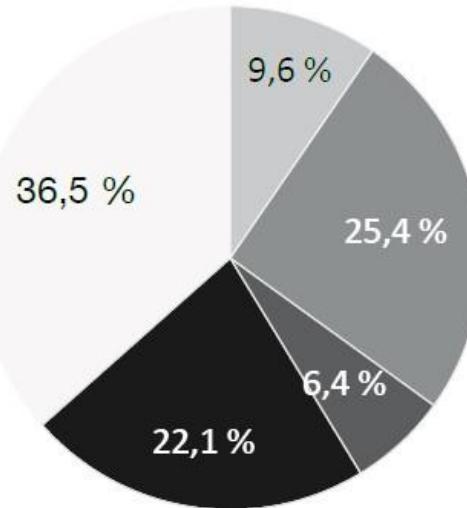
GRANDProto300 - Digital

338 tCO<sub>2</sub>e/yr



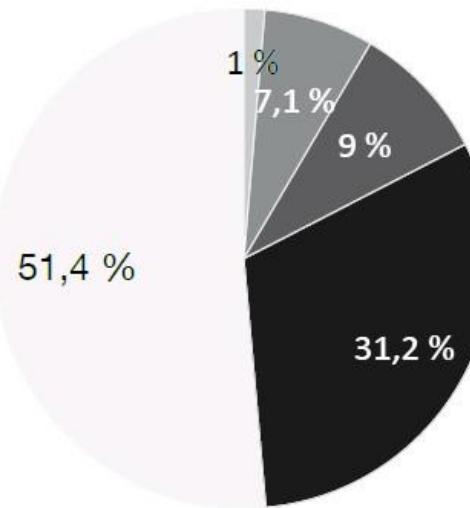
GRAND10k - Digital

422 tCO<sub>2</sub>e/yr



GRAND200k - Digital

6007 tCO<sub>2</sub>e/yr



- Devices
- Simulations
- Data analysis
- Data transfer
- Data storage

- Over- and underclocking with availability of renewables

- Over- and underclocking with availability of renewables
- Queuing system with availability of renewables

- Over- and underclocking with availability of renewables
- Queuing system with availability of renewables
- Generate your own electricity with renewables

- Over- and underclocking with availability of renewables
- Queuing system with availability of renewables
- Generate your own electricity with renewables
- Use the waste heat of your computations

- Over- and underclocking with availability of renewables
- Queuing system with availability of renewables
- Generate your own electricity with renewables
- Use the waste heat of your computations
- Move computations where:

- Over- and underclocking with availability of renewables
- Queuing system with availability of renewables
- Generate your own electricity with renewables
- Use the waste heat of your computations
- Move computations where:
  - Electricity is greener

- Over- and underclocking with availability of renewables
- Queuing system with availability of renewables
- Generate your own electricity with renewables
- Use the waste heat of your computations
- Move computations where:
  - Electricity is greener
  - Computing is more efficient (hyperscale data centers)

# The way forward

General ideas **computing**

# The way forward

General ideas **computing**

- evaluate data / processing demands when planning experiments

# The way forward

## General ideas **computing**

- evaluate data / processing demands when planning experiments
- reduce CO<sub>2eq</sub>-intensity of electricity (renewables, location, time)

# The way forward

## General ideas **computing**

- evaluate data / processing demands when planning experiments
- reduce CO<sub>2eq</sub>-intensity of electricity (renewables, location, time)
- improve utilization rate → reduce/share hardware → cloud

# The way forward

## General ideas **computing**

- evaluate data / processing demands when planning experiments
- reduce CO<sub>2eq</sub>-intensity of electricity (renewables, location, time)
- improve utilization rate → reduce/share hardware → cloud
- reduce data volume (lossy compression?)

# The way forward

## General ideas **computing**

- evaluate data / processing demands when planning experiments
- reduce CO<sub>2eq</sub>-intensity of electricity (renewables, location, time)
- improve utilization rate → reduce/share hardware → cloud
- reduce data volume (lossy compression?)
- share expensive simulations in scientific community

# The way forward

## General ideas **computing**

- evaluate data / processing demands when planning experiments
- reduce CO<sub>2eq</sub>-intensity of electricity (renewables, location, time)
- improve utilization rate → reduce/share hardware → cloud
- reduce data volume (lossy compression?)
- share expensive simulations in scientific community
- use computing waste heat

# The way forward

General ideas **computing**

- teach efficient programming

# The way forward

General ideas **computing**

- teach efficient programming
- monitor energy demand

# The way forward

## General ideas **computing**

- teach efficient programming
- monitor energy demand
- put CO<sub>2eq</sub>-estimates in acknowledgements of papers

# The way forward

## General ideas **computing**

- teach efficient programming
- monitor energy demand
- put CO<sub>2eq</sub>-estimates in acknowledgements of papers
- adapt codes to GPUs, install GPU systems

# The way forward

## General ideas **computing**

- teach efficient programming
- monitor energy demand
- put CO<sub>2eq</sub>-estimates in acknowledgements of papers
- adapt codes to GPUs, install GPU systems
- (use IT hardware longer)

# The way forward

## General ideas **computing**

- teach efficient programming
- monitor energy demand
- put CO<sub>2eq</sub>-estimates in acknowledgements of papers
- adapt codes to GPUs, install GPU systems
- (use IT hardware longer)

Most importantly:

- Judge which computations make sense

# The way forward

## General ideas **computing**

- teach efficient programming
- monitor energy demand
- put CO<sub>2eq</sub>-estimates in acknowledgements of papers
- adapt codes to GPUs, install GPU systems
- (use IT hardware longer)

Most importantly:

- Judge which computations make sense
- What is the code doing (purpose) that you are investing your time in

# The way forward

## General ideas **computing**

- teach efficient programming
- monitor energy demand
- put CO<sub>2eq</sub>-estimates in acknowledgements of papers
- adapt codes to GPUs, install GPU systems
- (use IT hardware longer)

Most importantly:

- Judge which computations make sense
- What is the code doing (purpose) that you are investing your time in
- What are you working on, how much infrastructure/energy is needed?

- Thank you for your attention
- Questions / Comments / Discussion

Jan Rybizki (IAB Weimar, formerly MPIA Heidelberg)  
[j.rybizki@iab-weimar.de](mailto:j.rybizki@iab-weimar.de)  
Twitter: @rybizki